



Datomic Cloud Documentation

- [Home](#) >
- [Operation](#) >
- [Howto](#)
- [Support](#)
- [Forum](#)

What is Datomic?

- [Data Model](#)
- [Architecture](#)
- [Supported Operations](#)
- [Programming with Data and EDN](#)

Local Dev and CI

Cloud Setup

- [Start a System](#)
- [Configure Access](#)
- [Get Connected](#)

Tutorial

- [Client API](#)
- [Assertion](#)
- [Read](#)
- [Accumulate](#)
- [Read Revisited](#)
- [Retract](#)
- [History](#)

Client API

Videos

- [AWS Setup](#)
- [Edn](#)
- [Datalog](#)
- [Datoms](#)
- [HTTP Direct](#)
- [CLI tools](#)

Schema

- [Defining Attributes](#)
- [Schema Reference](#)
- [Changing Schema](#)

- [Data Modeling](#)
- [Schema Limits](#)

Transactions

- [Processing Transactions](#)
- [Transaction Data Reference](#)
- [Transaction Functions](#)
- [ACID](#)
- [Client Synchronization](#)

Query and Pull

- [Executing Queries](#)
- [Query Data Reference](#)
- [Pull](#)
- [Index-pull](#)
- [Raw Index Access](#)

Time in Datomic

- [Log API](#)
- [Time Filters](#)

Ions

- [Ions Tutorial](#)
- [Ions Reference](#)
- [Monitoring Ions](#)

Analytics Support

- [Configuration](#)
- [Connecting](#)
- [Metaschema](#)
- [SQL CLI](#)
- [Troubleshooting](#)

Analytics Tools

- [Metabase](#)
- [R](#)
- [Python](#)
- [Jupyter](#)
- [Superset](#)
- [JDBC](#)
- [Other Tools](#)

Operation

- [Planning Your System](#)
- [Start a System](#)
- [AWS Account Setup](#)
- [Access Control](#)
- [CLI Tools](#)
- [Client Applications](#)
- [High Availability_\(HA\)](#)
- [Howto](#)

[Table of Contents](#)[Install the AWS CLI](#)[Manage AWS Access Keys for Datomic](#)[Install Clojure CLI](#)[Install ion-dev Tools](#)[Install ion Library](#)[Delete Ion lambdas](#)[Control Shell Scripts](#)[Find Datomic System Name](#)[Find Datomic Nodes](#)[Find Compute Group Name](#)[Check System Topology](#)[Find Datomic Application Name](#)[Find System S3 Bucket](#)[Find Ion Code Bucket](#)[Upgrade Base Schema](#)[Update a CloudFormation Parameter](#)[Convert from Solo to Production](#)

- [Query Groups](#)
- [Monitoring](#)
- [Upgrading](#)
- [Scaling](#)
- [Deleting](#)
- [Splitting Stacks](#)

[Tech Notes](#)

- [Turning Off Unused Resources](#)
- [Reserved Instances](#)
- [Lambda Provisioned Concurrency](#)

[Best Practices](#)

[Troubleshooting](#)

[FAQ](#)

[Examples](#)

[Releases](#)

[Glossary](#)

Hide All Examples

How To

Table of Contents

- [Install the AWS CLI](#)
- [Manage AWS Access Keys for Datomic](#)
- [Install Clojure CLI](#)
- [Install ion-dev Tools](#)
- [Install ion Library](#)
- [Delete Ion lambdas](#)
- [Control Shell Scripts](#)
- [Find Datomic System Name](#)
- [Find Datomic Nodes](#)
- [Find Compute Group Name](#)
- [Check System Topology](#)
- [Find Datomic Application Name](#)
- [Find System S3 Bucket](#)
- [Find Ion Code Bucket](#)
- [Upgrade Base Schema](#)
- [Update a CloudFormation Parameter](#)
- [Convert from Solo to Production](#)

Install the AWS CLI

[Install the AWS Command Line Interface](#). Make sure you have version 1.11.170 or later. If you are unfamiliar with managing Python environments, first try using the bundled [AWS CLI installer](#). You can check your version using:

```
aws --version
```

NOTE: Running the `datomic-socks-proxy` script with an earlier version of the AWS CLI will result in a [generic error message](#) and failure of the script.

If you require multiple versions of the AWS CLI to be installed, then please investigate [using Python's virtual environments](#), or the third-party [virtualenv solution](#).

Manage AWS Access Keys for Datomic

The Datomic [CLI tools](#) and [ion deployment tools](#) require that you have a set of AWS access key with permissions to access Datomic. To create these keys, you must have an [authorized IAM user](#). With this user you can then:

1. [Create access keys](#)
2. [Add access keys to your environment](#)

Datomic supports the use of [named profiles](#) as a credentials source. Additional information about IAM Users and access keys can be found in the [AWS Security Credentials documentation](#).

Install Clojure CLI

Datomic's tools use the [Clojure CLI](#), version 1.10.1.478 or later. You can use `clojure -sdescribe` to check your version of the Clojure CLI.

The Clojure [Getting Started](#) page has instructions for installing and upgrading Clojure.

Install ion-dev Tools

The Datomic ion-dev tools should be installed via an alias in your [user deps.edn file](#), which is normally located at `$HOME/.clojure/deps.edn`.

To install the tools, add the `datomic-cloud` maven repo under your `:mvn/repos` key:

```
"datomic-cloud" {:url "s3://datomic-releases-1fc2183a/maven/releases"}
```

Then add an `:ion-dev` entry under your `:aliases` key with the [latest version](#) of ion-dev:

```
:ion-dev
{:deps {com.datomic/ion-dev {:mvn/version "0.9.247"}}
 :main-opts ["-m" "datomic.ion.dev"]}
```



The ions tutorial includes a [complete .clojure/deps.edn example](#).

Note that `ion-dev` configures logging to stderr via [slf4j-simple](#). This is probably adequate for most scenarios, but you of course have the [full power of SLF4J](#) at your disposal.

Install ion Library

The Datomic ion library should be installed in the `deps.edn` for each ion project:

Add the `datomic-cloud` maven repo under your `:mvn/repos` key:

```
"datomic-cloud" {:url "s3://datomic-releases-1fc2183a/maven/releases"}
```

Then add the [latest version](#) of `ion-dev` to your `:deps`:

```
com.datomic/ion {:mvn/version "0.9.35"}
```

Delete Ion lambdas

You do **not** need to delete the AWS Lambdas managed by Datomic Ions in order to control cost. AWS Lambda pricing is entirely usage-based, so if you do not invoke Lambdas, they cost nothing. If you want to delete these Lambdas anyway, you can [push](#) and [deploy](#) an application with an empty `lambdas` map in [ion-config.edn](#).

Control Shell Scripts

Some Datomic CLI tools (e.g. the [client access](#) script) continue to run in the foreground once you launch them. For interactive use, the easiest way to manage such tools is simply to kill them with Ctrl-C when they are no longer needed.

If you are building automation around scripts, you can of course use all the ordinary Unix facilities, e.g.

- You can use `kill [script-name]` to kill a script from another terminal.
- You can use e.g. `nohup [script-command] [script-args] &` to launch a script in the background, directing its output to `nohup.out`.

Find Datomic System Name

Every Datomic system resides in a single AWS region, and has a unique system name within that region. The system name is the Stack Name used to [launch the Datomic storage stack](#).

If you did not [start Datomic](#) yourself, you can find the available systems in a region two ways:

- Utilize the [datomic cloud list-systems](#) command.

or

- Browse to [CloudFormation console](#).

Find Datomic Nodes

You can list the names of all running Datomic nodes using:

- Utilize the [datomic system list-instances <system>](#) command, replacing `<system>` with your system name.

or

- the following [AWS CLI](#) command, replacing `[region]` with the region you want to query:

```
aws ec2 describe-instances --region [region] --filters "Name=tag-key,Values=datomic:tx-group" "Name=tag-value,Values=datomic:tx-group"
```

Find Compute Group Name

- If you start a Datomic system with [split stacks](#), then it has one or more compute groups. A compute group's name is the name of the CloudFormation stack you used to [create the group](#).
- If you start a system [using the AWS Marketplace template](#), Datomic generates a unique name for your compute stack based on your [system name](#), of the form \$(SystemName)-Compute-\$(GeneratedId).

You can:

- Utilize the [datomic system list-instances <system>](#) command, replacing <system> with your system name.

or

- browse all your CloudFormation stacks in the [AWS CloudFormation console](#).

Check System Topology

All Datomic systems include one or more CloudFormation stacks. To check whether a system uses the [Solo or Production topology](#) either:

- Utilize the [datomic cloud list-systems](#) command.

or

- Find your system's [primary compute group](#) in the CloudFormation console.
- If the CloudFormation template has a DatomicProductionCompute key in its Outputs, the system is running the Production topology. If the key is absent, the system is running Solo.

Find Datomic Application Name

Every Datomic [compute group](#) has an associated [application](#) for ion deployments. You specify the application name when you [create a compute group](#).

If you did not start the compute group, you can find the application name by browsing to your compute group in the [AWS CloudFormation console](#) and looking for the ApplicationName value in the Parameters tab.

Find System S3 Bucket

Every Datomic system has its own S3 bucket for configuration and data storage.

You can find this S3 bucket by browsing to the Outputs tab of your [storage stack](#) in the [AWS Console](#) and looking for the the S3DatomicArn output.

Find Ion Code Bucket

All Datomic systems in a region share a common S3 bucket for [ion](#) deployment. This bucket has a datomic:code AWS tag.

You can find this bucket by creating a [resource group](#) in the AWS Console:

- make sure you are working in the region you want to query
- create a tag based resource group
- select a resource type of AWS::S3::Bucket
- search for a tag named datomic:code
- give you group a name, e.g. DatomicCodeBucket
- create your group

Upgrade Base Schema

Every Datomic database starts with a base schema, whose identifiers are prefixed with `:db`. When the Datomic team enhances the base schema, all databases created with new versions of Datomic get the enhancements automatically.

Existing databases do **not** automatically incorporate enhancements to the base schema. In order to upgrade existing databases to use new base schema, you must first upgrade **all** compute groups (primary and query) to the minimum version of Datomic required for the base schema features you want:

Feature	Minimum Datomic Version	Released
tuples	480-8770	June 27, 2019
attribute predicates	480-8770	June 27, 2019
entity specs	480-8770	June 27, 2019
db.type/symbol	480-8770	June 27, 2019

Once you have upgraded all compute groups, you can upgrade an existing database to use the latest base schema by passing the `:upgrade-schema` action to `administer-system`. For example:

```
(d/administer-system client {:db-name "movies"
                           :action :upgrade-schema})
```

Upgrading a schema is an idempotent operation, so it will not harm your data to call it more than once. That said, you should treat `administer-system` as a potentially expensive operation. Call it only when you actually need it, not before every call to `connect`.

Once you have used a base schema feature in a particular database, do **not** run any compute group for that system on versions of Datomic older than the features used. (See table above.)

Update a CloudFormation Parameter

Datomic's CloudFormation parameter settings are carefully planned to adhere to AWS best practices and provide a robust and performant system. You should change CloudFormation parameters only:

- when the Datomic documentation specifically instructs you to do so
- when the Datomic support team specifically instructs you to do so

To update a CloudFormation parameter:

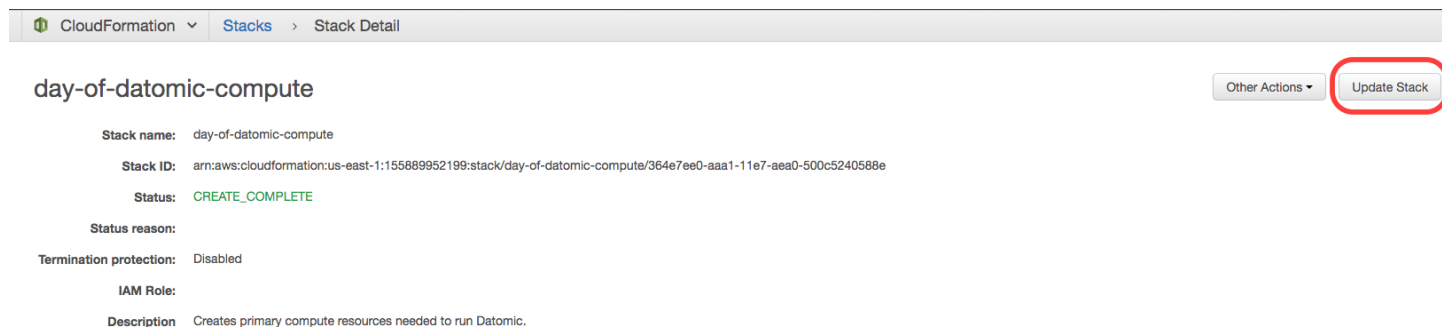
- Select your stack via the checkbox or radio button.
- Click the "Actions" button and select "Update Stack"
- On the "Select Template" screen, choose "Use current template"
- On the "Specify Details" screen, change your parameter(s).
- On the "Options" screen, leave all options unchanged.
- On the "Review" screen, click the checkbox stating "I acknowledge that AWS CloudFormation might create IAM resources with custom names."

Convert from Solo to Production

To update your system from Solo to Production [topology](#):

- make sure you have [raised](#) your i3.large instance limit
- click the name of your solo stack in

[CloudFormation console](#). On the stack page, Click "Update Stack":



On the following page, select "Specify an Amazon S3 template URL:" and enter the CloudFormation template URL for the latest production template you wish to upgrade to (see [Release page](#) for the latest production template) and click "Next".

- On the "Specify Details" screen, leave all options unchanged.
- On the "Options" screen, leave all options unchanged.
- On the "Review" screen, click the checkbox stating "I acknowledge that AWS CloudFormation might create IAM resources with custom names."

Wait for the template to report UPDATE_COMPLETE. This can take five or more minutes. You can refresh the CloudFormation dashboard to see progress. In the [EC2 Dashboard](#) you'll note that two i3.large instances with your stack name are running and a t2.small instance has been terminated. Your production stack is now ready for use.

Copyright © Cognitect, Inc

Datomic® and the Datomic logo are registered trademarks of Cognitect, Inc