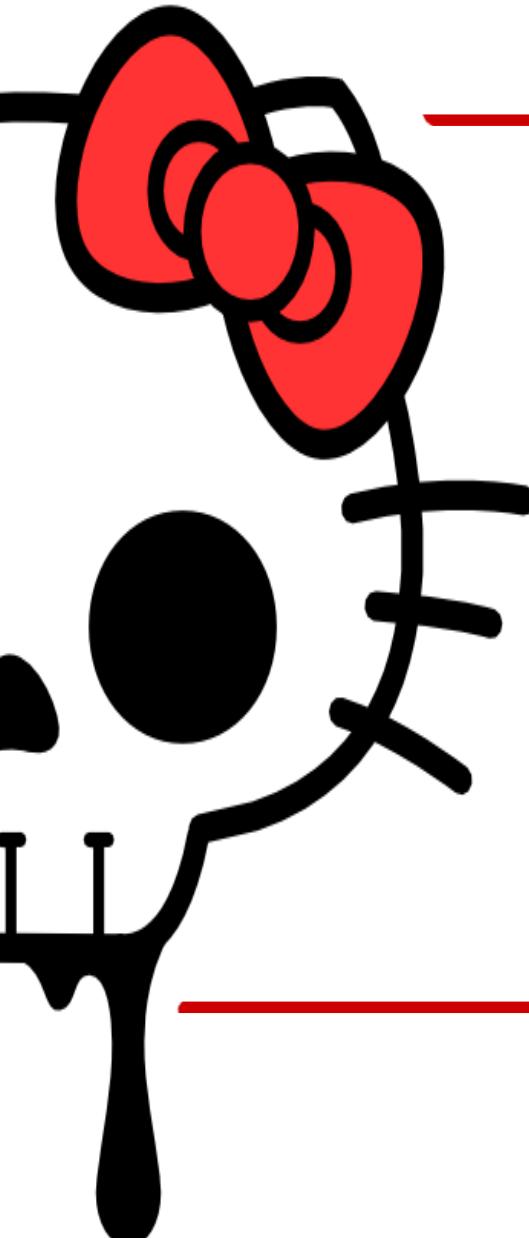


Glibc Heap Analysis in Linux Systems with Radare2



`(:)\> forget what we made _`

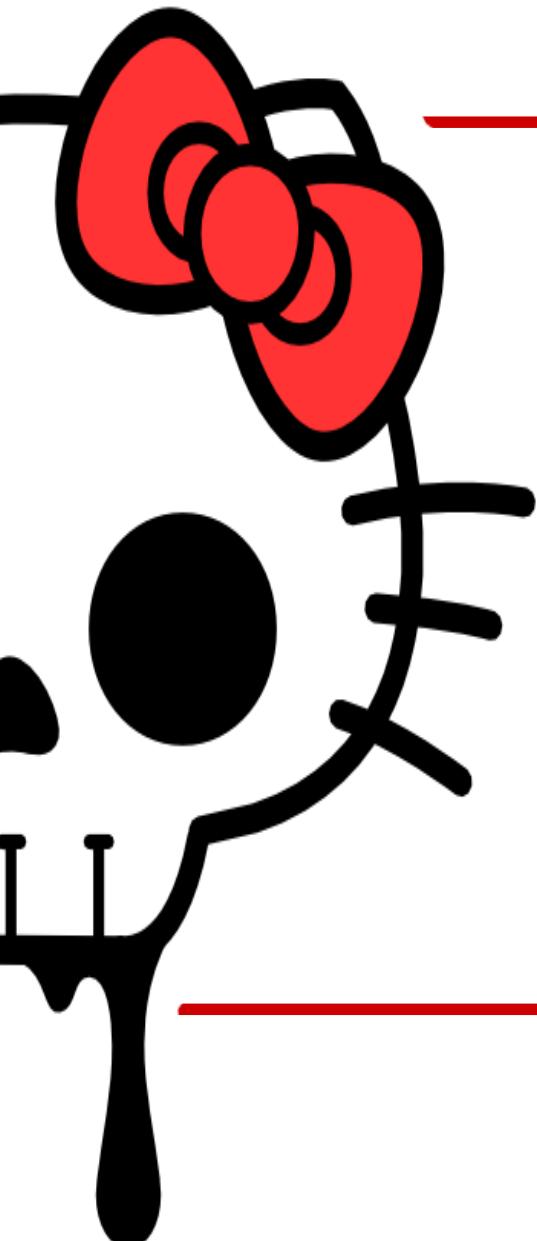
Introduction



```
n4x0r@r2con:~/r2_heap$ whoami
```

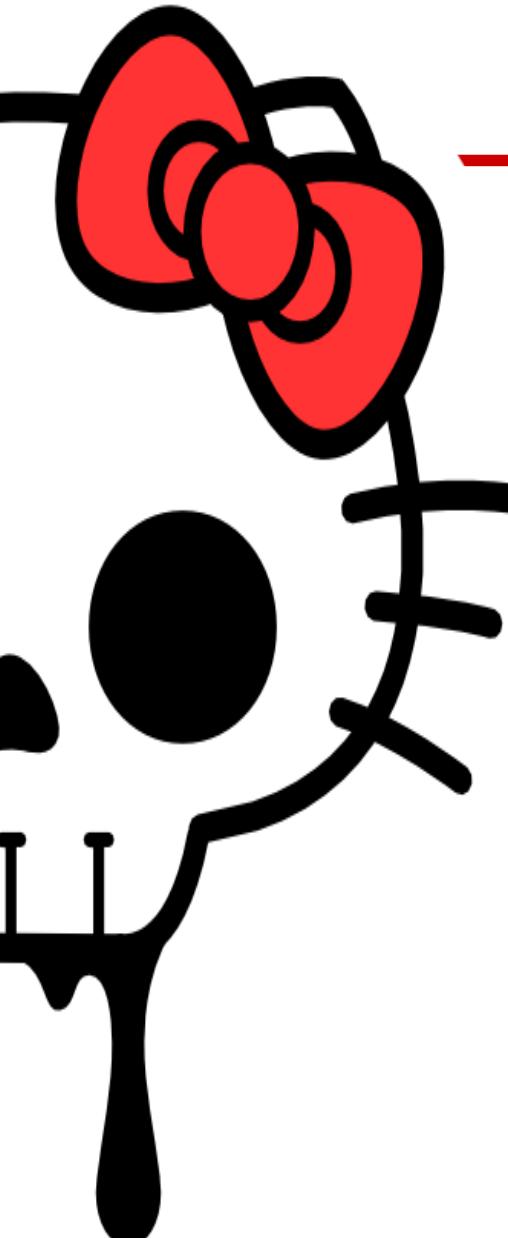
- CTF player at am3s1a_team.
- Computer Science Undergraduate at The University of Liverpool, UK.
- Mainly focused on binary Exploitation and Reverse Engineering.

Talk Layout



1. Contributors
2. Motivation
3. Glibc malloc Heap Essentials
4. Dmh command family
5. Conclusion
6. Questions

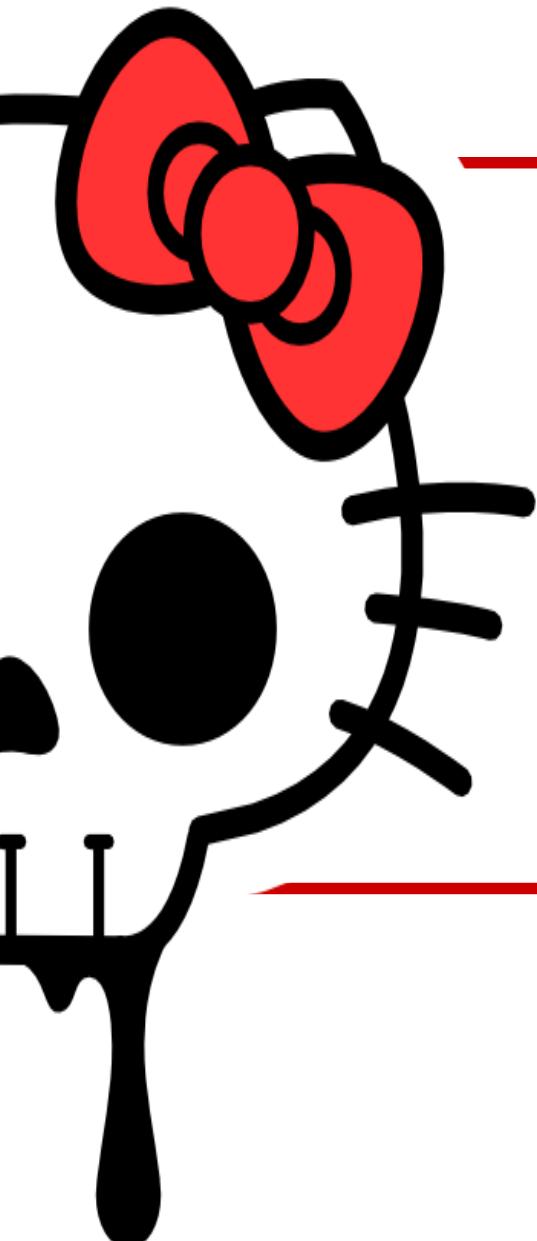
Contributors



```
n4x0r@r2con:~/r2_heap$ ls -l
```

- -rwxr--r-- 1 @javierprtd @amn3s1a_team 373 Jul 9 23:50 Javier_Partido
- -rwxr--r-- 1 @n4x0r_ @amn3s1a_team 373 Jul 9 23:50 Ignacio_Sanmillan

Motivation



- Mainly focused on Exploitation and Reverse Engineering in CTF competitions.
- Seeking for a tool to analyze glibc malloc Heap.
- Found lack of quality features in known debuggers to analyze the Heap efficiently.
- GDB plugins did not meet expectations.

- Outcome: To Implement something on our own.
- Problem: Choosing framework to implement it.
- Solution: Radare2

Glibc malloc Heap Essentials

Glibc malloc Heap Essentials

- Heap Overview

- Heap creation

- Heap structures
 - + Malloc_chunk
 - + Malloc_state
 - + Heap_Info

- Intro to glibc malloc algorithm

Heap Overview



- **What is the heap?**

Region of memory created at runtime intended to allocate variables which size is not deterministic at compile time.

- **2 functions which make up the foundations of heap manipulation. These functions are:**

Malloc:

- Prototype: **void* malloc(size_t size)**
- Functionality:
 - First time malloc is called in thread:
 - Allocates a reasonable amount of memory,
 - Creates a heap segment or equivalent,
 - Returns the caller a pointer to a memory region within it of suitable for demand.
 - Not the first call to malloc in thread:
 - malloc will just return a pointer to a region within the current heap segment (or equivalent) of suitable size for the caller demand.

Free:

- **Prototype: void free(void *ptr)**
- **Functionality:**
 - Set a memory region previously returned from malloc as not in use.

- **A little history about allocators**

The heap segment has to be created at runtime by an algorithm.

Algorithms that allocate memory dynamically and administrate it are denominated allocators.

- **There are various allocators that use malloc, free ...etc, however even though their functionality resembles similar, their implementation is quite different.**

Some known allocators are:

- **Dmalloc: General purpose allocator**
- **Ptmalloc: (ptmalloc2) Glibc**
- **Jemalloc: FreeBSD and Firefox**
- **Tcmalloc: Google Chrome**
- **Magazine malloc: IOS/OSX**

This talk is based on the glibc malloc algorithm (glibc version 2.24) since is the most wide spreaded within linux distributions.

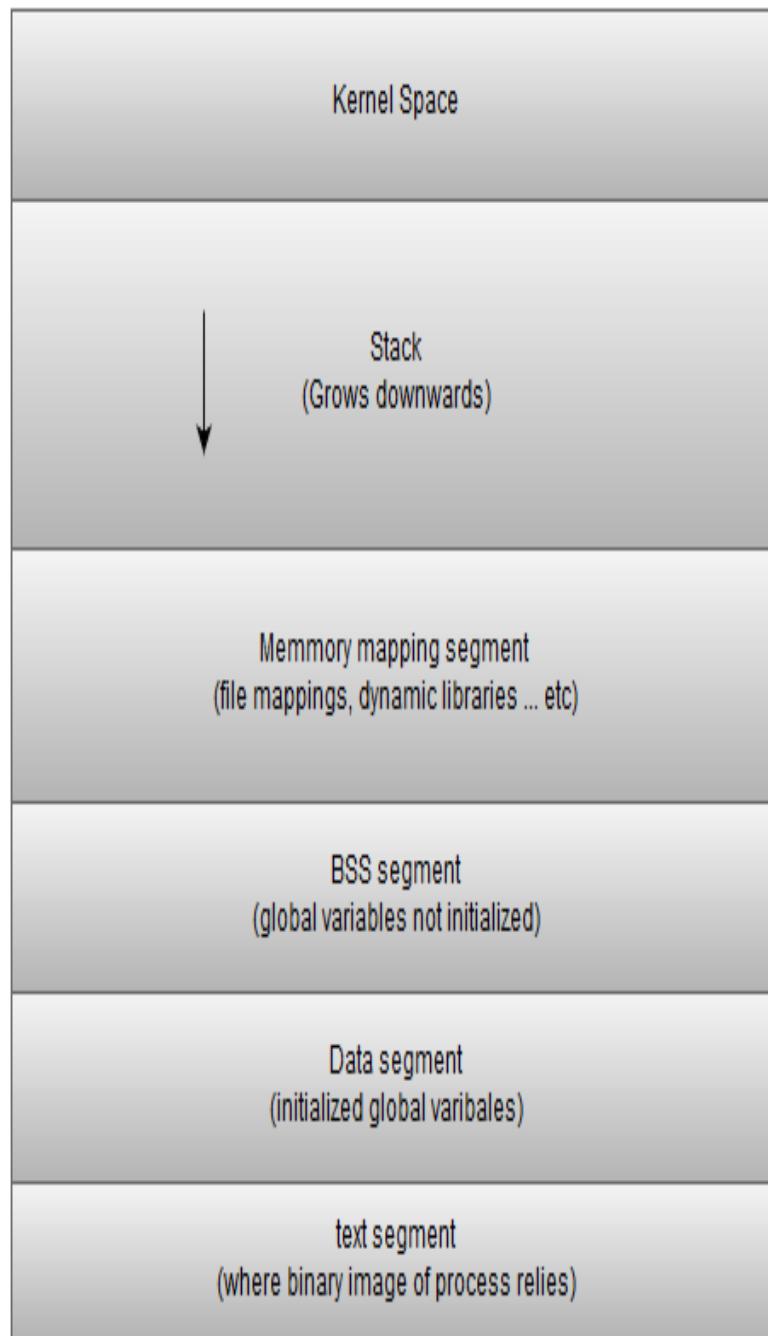
- **A little history about glibc malloc:**
 - Based on the ptmalloc allocator (pthreads malloc) version 2, which was originally derived from dlmalloc allocator.
 - Major differences between ptmalloc2 and dlmalloc:
 - ptmalloc2 supports heaps on different threads as well as multiple heaps in child threads.
 - The actual implementation of glibc malloc may vary from ptmalloc2 implementation due to different changes in glibc malloc from its ptmalloc2 integration.

- **Some definitions:**
 - **End_data:** pointer which sets the end of the data segment.
 - **start_brk:** pointer which denotes the beginning of the heap segment.
 - **Brk / program break:** pointer that denotes the end of the heap segment.
- Allocators communicate with the kernel in order to retrieve memory from the OS via systemcalls. In glibc malloc these system calls are the following:

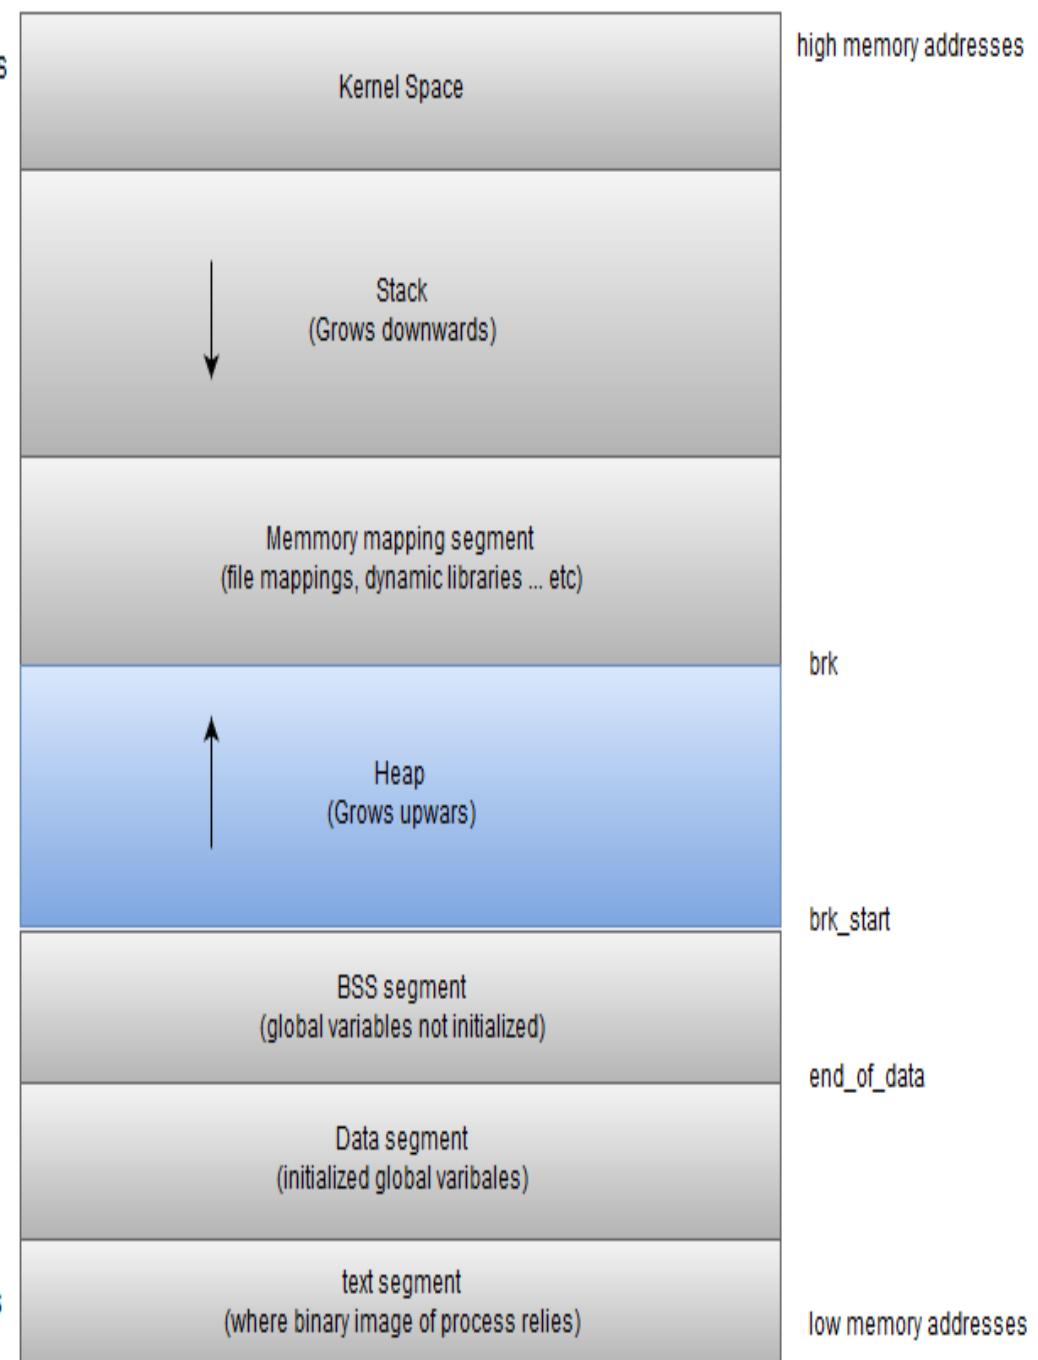
- **The brk() systemcall:**
 - Prototype: int brk(void *addr);
 - **Initializes:**
 - **end_data to the value specified by addr.**
 - **start_brk to the next available segment after the data segment.**
 - **brk = start_brk.**
 - **In order to expand or shrink the dimensions of the heap segment, an additional glibc function is used. This function is sbrk().**

- **The sbrk() function:**
 - Prototype: **void *sbrk(intptr_t increment);**
 - **wrapper function around brk(),**
 - **increments the program break by increment bytes, therefore giving dimensions to the heap segment.**
 - **Calling sbrk() with an argument of 0 can be used to find the current location of the program break.**

first call to brk() by malloc

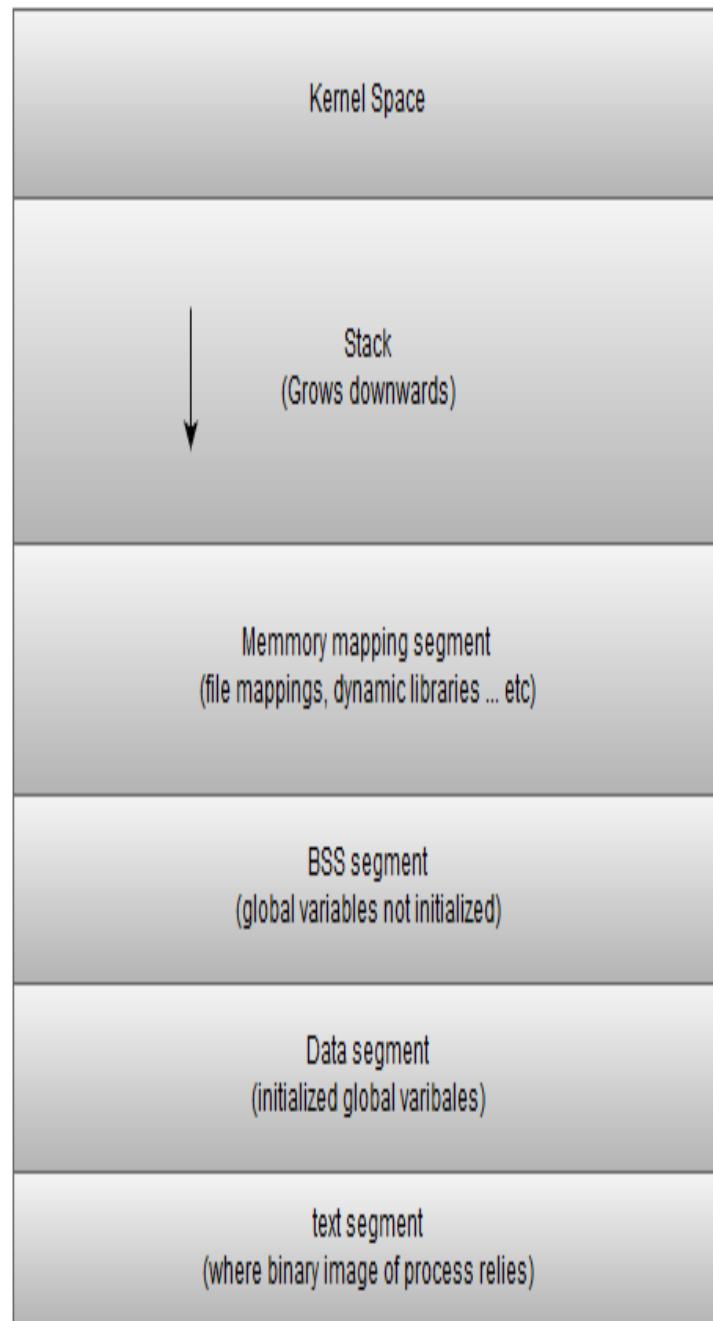


first call to sbrk() by malloc



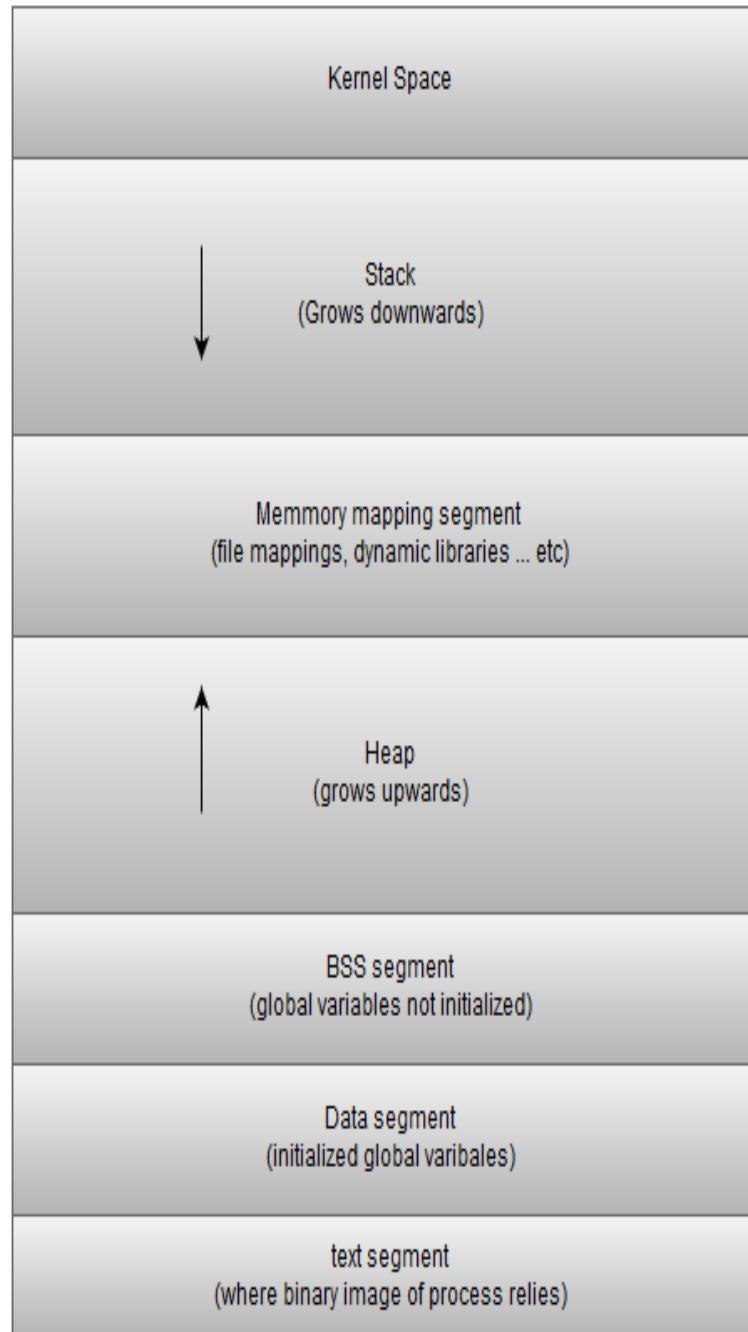
before calling malloc for the first time

After calling malloc for the first time



high memory addresses

low memory addresses



high memory addresses

brk

start_brk

end_of_data

low memory addresses

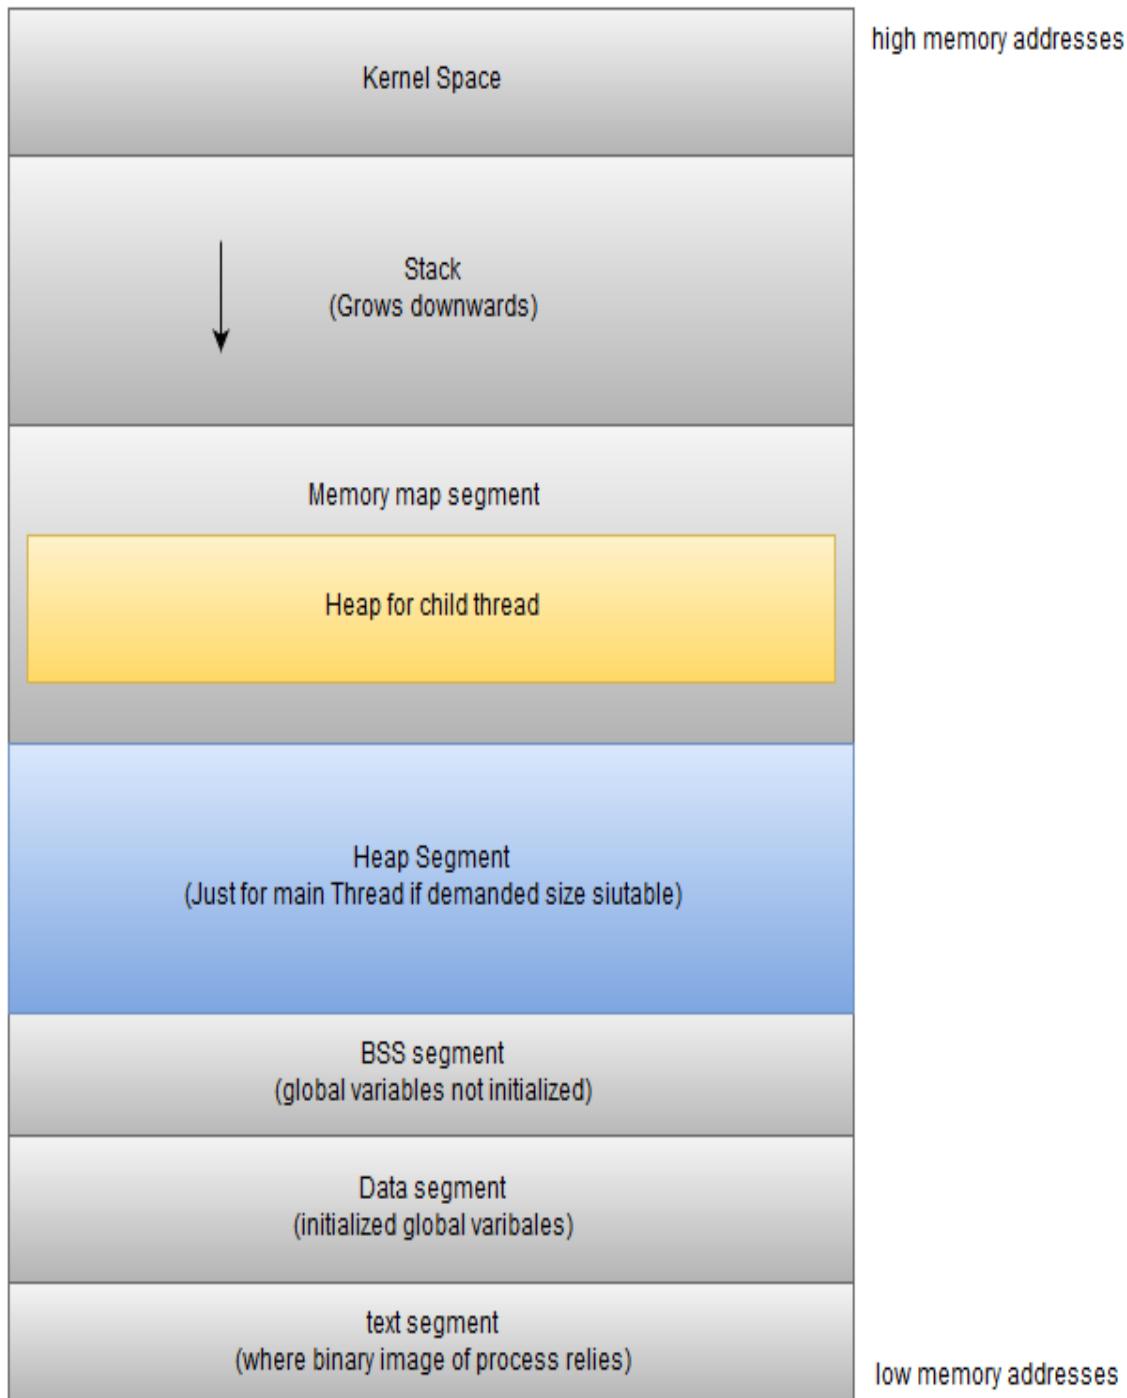


Brk() / sbrk() only create the heap segment for the main thread!.

- **The mmap() systemcall:**

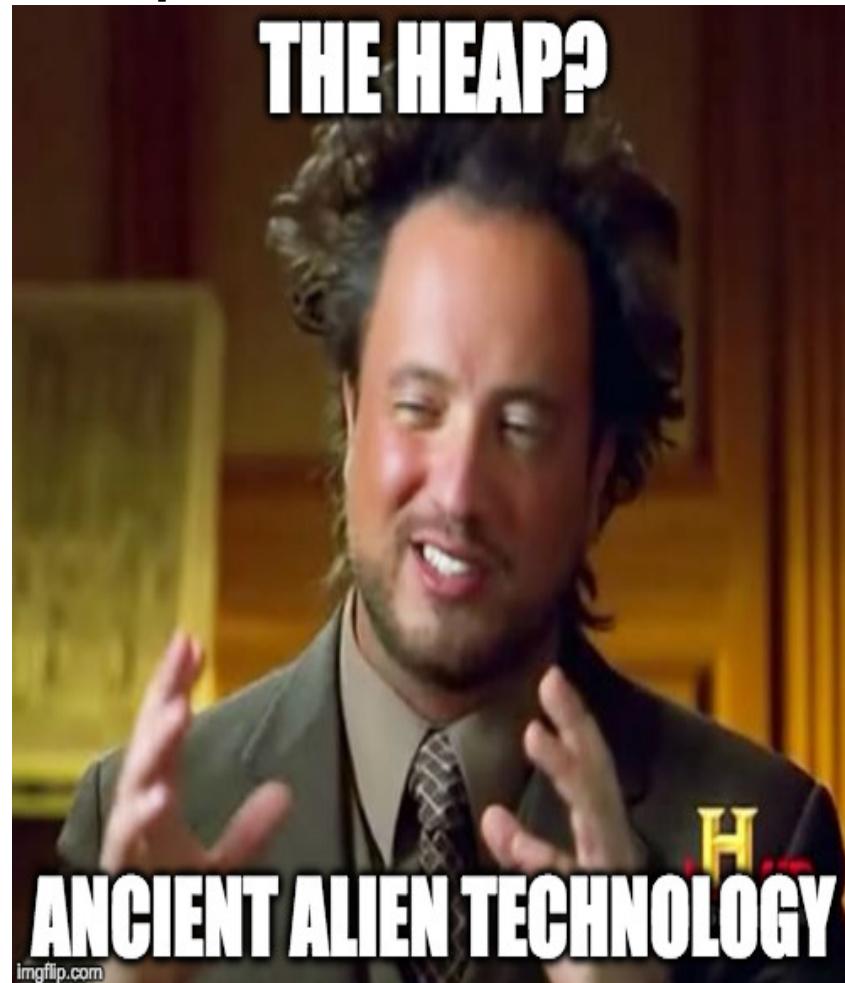
- **Prototype: void * mmap (void *address, size_t length, int protect, int flags, int filedes, off_t offset)**
- **Functionality:**
 - Mmap() will be called in 2 different scenarios:
 - 1) Demanded size for a malloc call in the main thread exceeded the macro DEFAULT_MMAP_THRESHOLD which has a value of 128 * 1024
 - 2) Call to malloc proceeded from a non main thread.

- **How does mmap() work?:**
 - **treats child “heap segment” as private anonymous allocations.**
 - **Allocates a region of memory in the memory mappings segment filled with 0s, unlinked from any file and being private for each thread.**
 - **Mmap si capable to do this with the following flags:**
 - **MAP_ANON**
 - **MAP_PRIVATE**

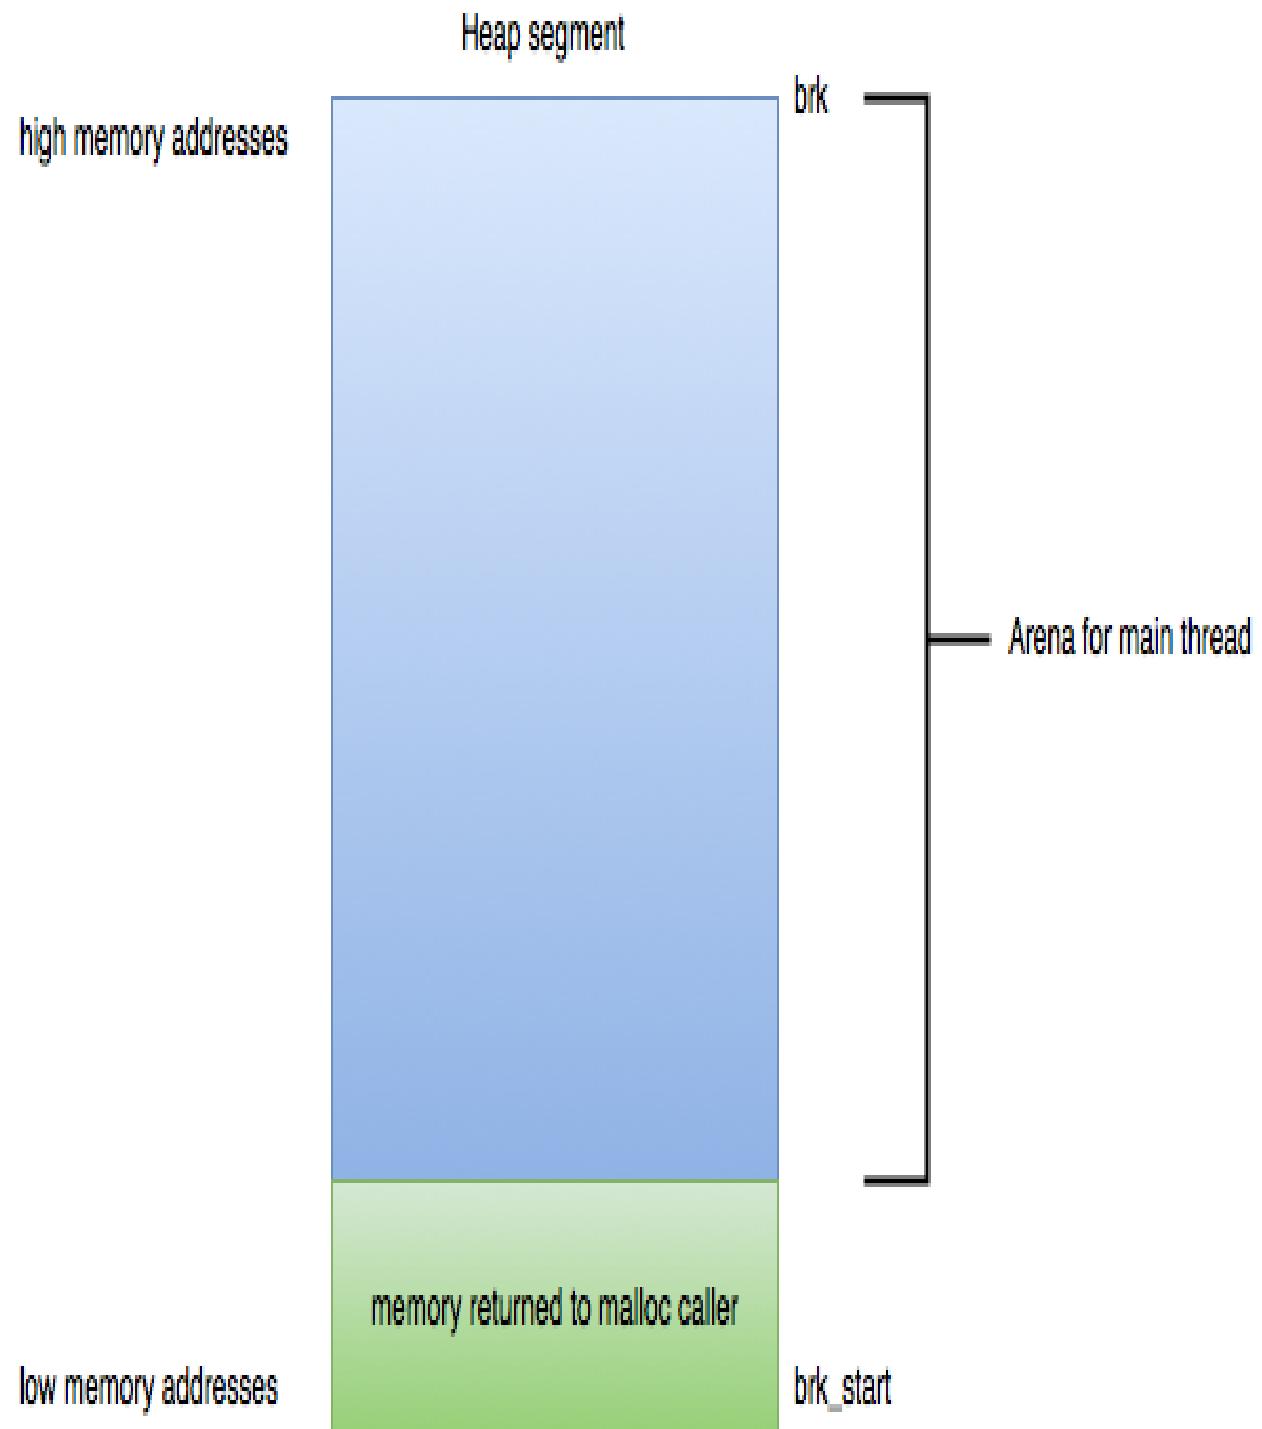


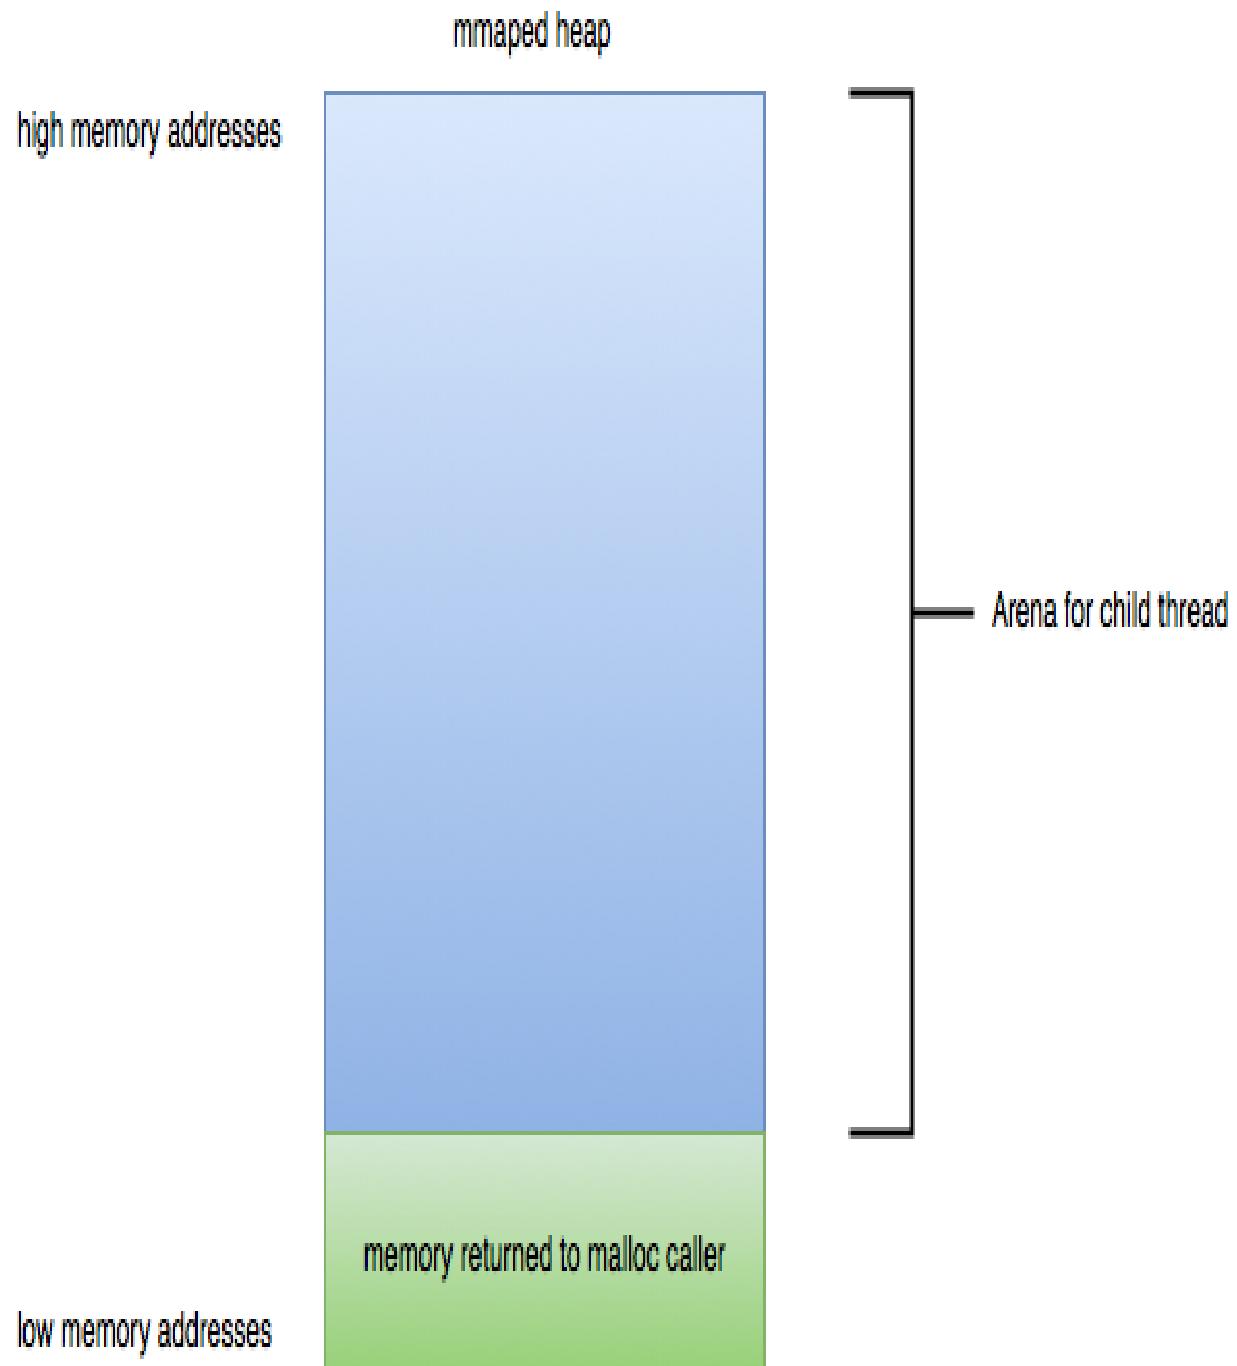
Heap Organization

- Before digging any further into glibc malloc heap internas, first we have to understand a few concepts:



- **Arena:**
 - Product of optimization procedure to avoid unnecessary interaction with the kernel while requesting memory for dynamic allocation.
 - An arena is all the available memory left that was retrieved from the OS to serve all potential malloc calls demands for a particular thread.





- **There is a limit to the number of arenas that a concurrent application may support.**
- **This limit is determined based on the number of cores the system has.**
 - **limit = (32bit) ? number of cores * 2 : number of cores * 8.**
 - **This limit can also be controlled using the MALLOC_ARENA_MAX environment variable.**

- **Heap**

Nothing but a contiguous region of memory that is composed of pieces called chunks.

Each heap belongs to exactly one arena.

- **Heap != Heap segment.**

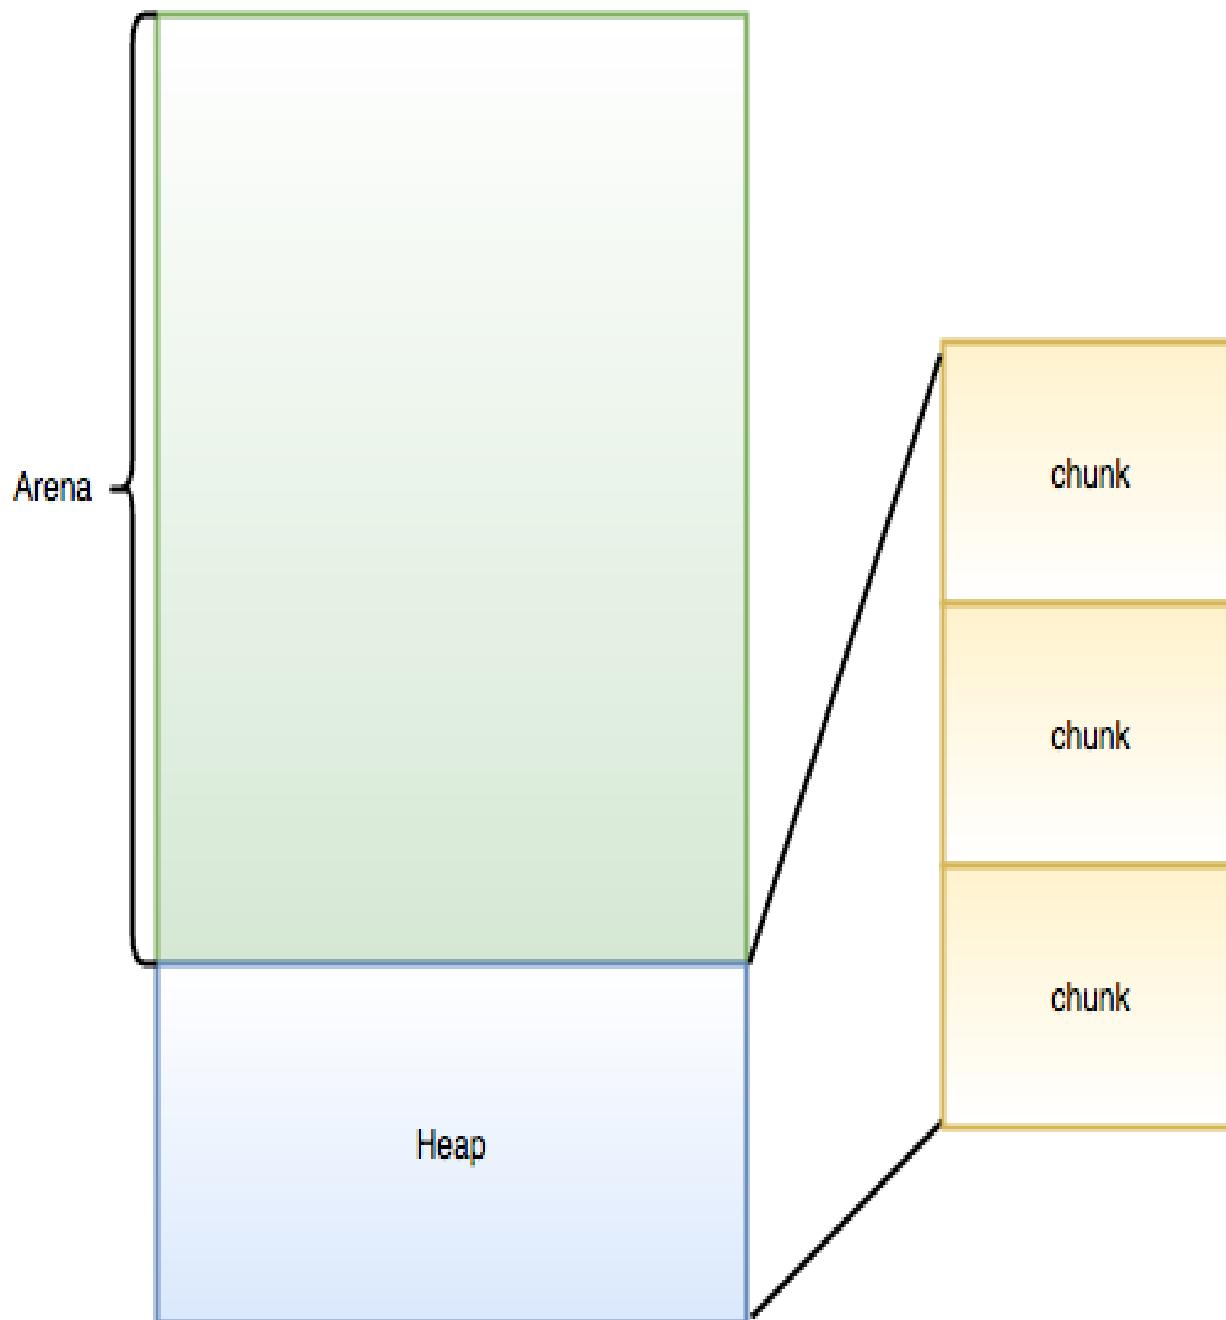
- **Chunk**

Smallest unit of memory administration in dynamic allocation.

When malloc is called, it returns a pointer to a chunk with suitable size for the particular call.

Each chunk exists in one heap and belongs to one arena.

Conceptual model of Arena, Heap and Chunk



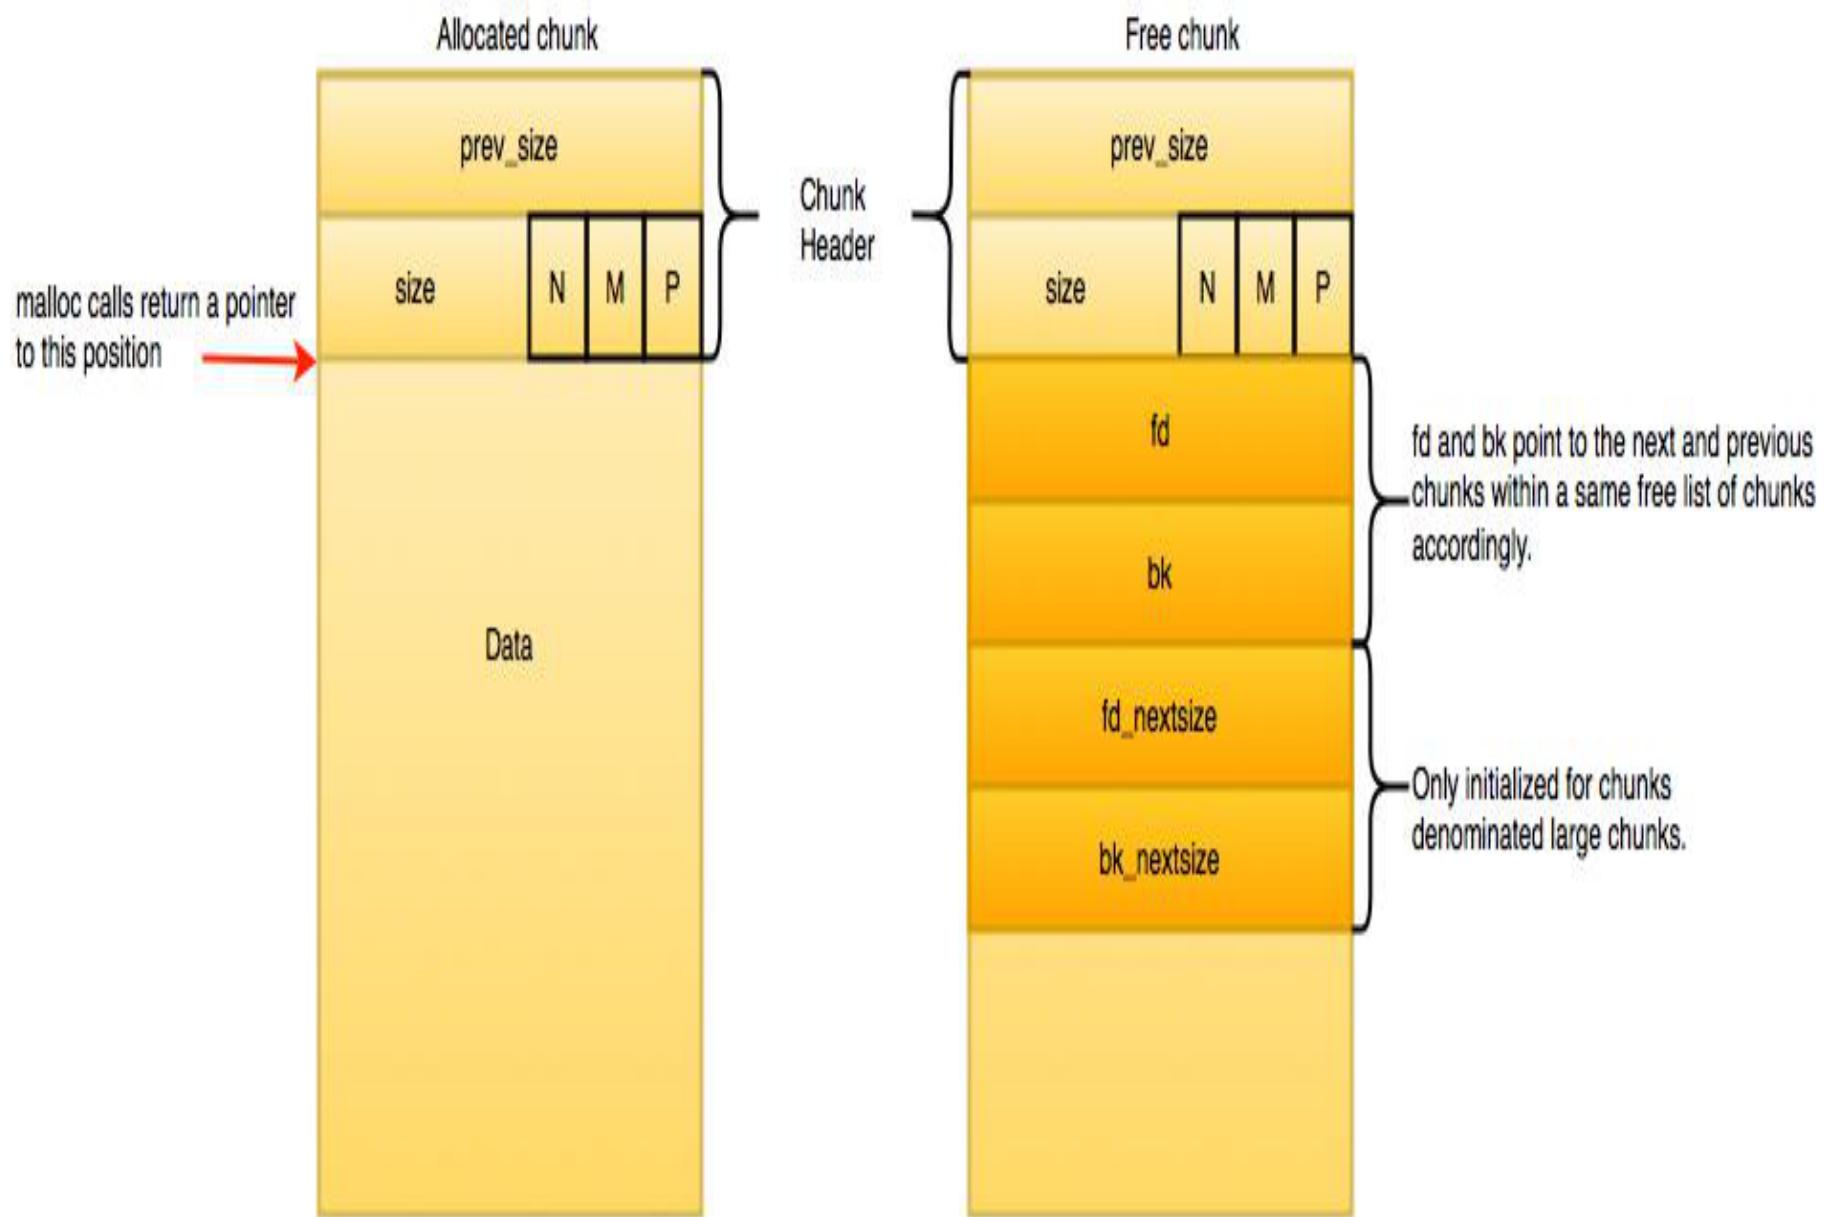
Heap Structures

- The previous concepts are implemented in the form of structures. A few of the most relevant structures in glibc malloc are the following:
 - **malloc_chunk** : all the elements that makeup a chunk.
 - **malloc_info**: all elements that interact with the heap or heaps for a particular child thread.
 - **malloc_state**: all elements that compose an arena.

malloc_chunk:

```
struct malloc_chunk {  
  
    INTERNAL_SIZE_T prev_size; /* Size of previous chunk (if free). */  
    INTERNAL_SIZE_T size;      /* Size in bytes, including header. */  
  
    struct malloc_chunk* fd;    /* double links -- used only if free. */  
    struct malloc_chunk* bk;  
  
    /* Only used for large blocks: pointer to next larger size. */  
    struct malloc_chunk* fd_nextsize; /* double links -- used only if free. */  
    struct malloc_chunk* bk_nextsize;  
};
```

- Chunks can either be free or allocated.
- Their structure change accordingly
 - prev_size field is only initialized if chunk is free.
 - When chunks are freed, they get inserted in linked list of free chunks.
 - Fd, bk are then initialized.
 - Link a particular chunk to others within a same linked list.
 - (more about this in coming slides).
- 3 LSB of size field represent chunk flags, these are:
 - PREV_INUSE (P) : bit set when previous chunk in heap is allocated.
 - IS_MAPPED (M) : bit set when chunk is being mmap'd.
 - NON_MAIN_arena (N): bit set when chunk does not belong to heap segment



malloc_state:

- Two possible types of instances for this structure:
 - **main_arena**: instance for the main thread arena. Held in glibc data segment as a global variable.
 - **thread_arena**: instance for any other arena which does not belong to the main thread.

Heap Structures



```
typedef struct malloc_state{  
  
    mutex_t mutex;           // serialize access  
    int flags;              // flags  
    mfastbinptr fastbinsY[NFASTBINS]; //array of single linked list of free chunks  
  
    mchunkptr top;          // base of top chunk  
    mchunkptr last_remainder; // base of remainder chunk  
  
    mchunkptr bins[NBINS * 2 - 2]; // array of double linked list of free chunks  
    unsigned int binmap[BINMAPSIZE]; // bitmap of bins  
    struct malloc_state *next;      // arena single linked list  
    struct malloc_state *next_free; // free arena linked list  
  
    INTERNAL_SIZE_T system_mem;    // current memory allocated for this arena  
    INTERNAL_SIZE_T max_system_mem;  
};
```

```
mfastbinptr fastbinsY[NFASTBINS]; //array of single linked list of  
free chunks
```

```
mchunkptr top; // base of top chunk  
mchunkptr last_remainder; // base of remainder chunk  
  
mchunkptr bins[NBINS * 2 - 2]; // array of double linked list of  
free chunks
```

Free Lists:

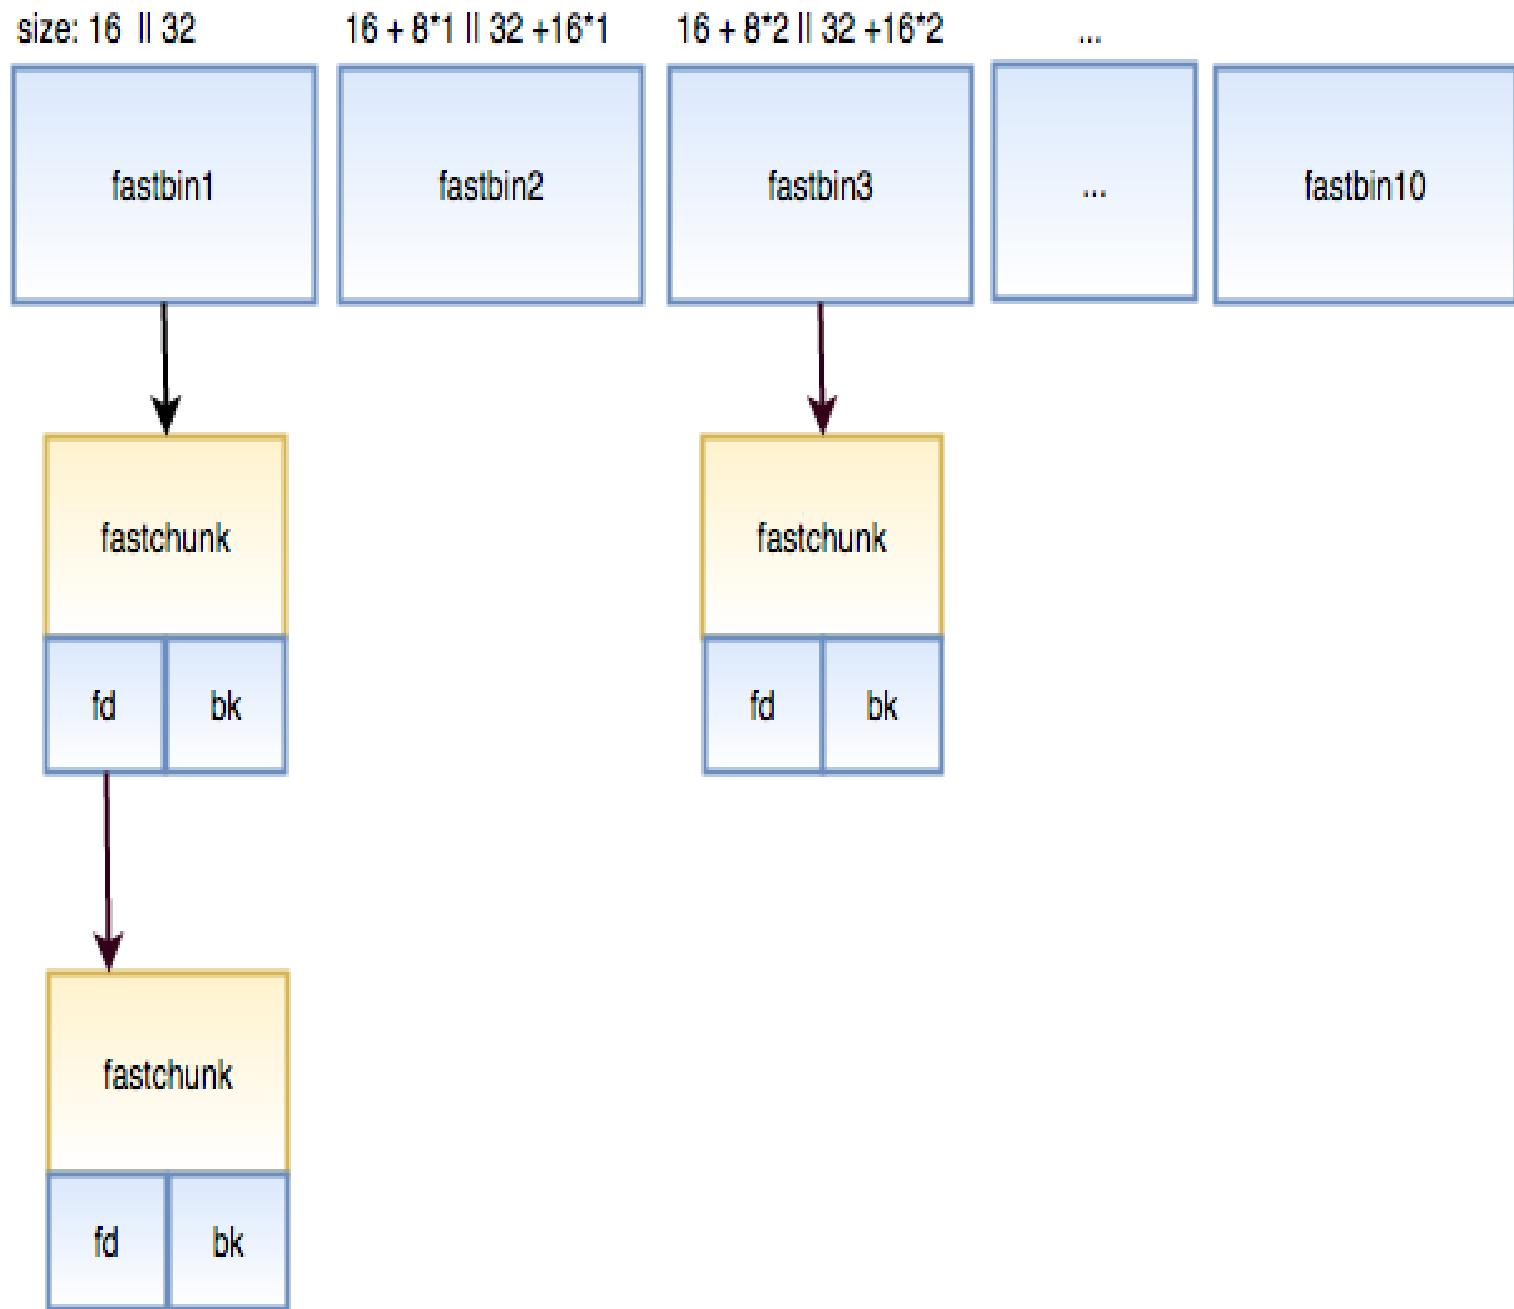
What happens when free() is called?:

Chunk is marked as free, and its reference will be placed in one of these two data structures:

- **fastbinsY**
- **bins**
- **The purpose of these two data structures is to recycle already allocated chunks by malloc to serve future malloc calls.**
- **Less kernel interaction == More efficiency**

FastbinY:

- **Array of LIFO single linked lists holding recently freed chunks denominated fastchunks.**
- **Fastchunks may have sizes rangin from:**
 - **Fastchunk max size:**
 - 80 bytes - 32bit,
 - 160 bytes - 64bit.
 - **Fastchunk min size:**
 - 16 bytes - 32bit.
 - 32 bytes - 64bit.
- **List that hold fastchunks is denominated a Fastbin.**
- **Fastbin number: 10.**
- **Very fast for allocation and deallocation.**
- **Each fastbin hold chunks of the same size**
- **Chunk's sizes in different fastbins are spaced by:**
 - **8 bytes in 32bit**
 - **16 bytes in 64bit.**

.fastbinY

Bins:

- **Each index in array is a FIFO double linked list of free chunks known as bin.**
- **Bin number: 128**
- **A particular bin holds free chunks of specific size/sizes depending on index.**

- **Bins for chunk sizes < 512 || 1024 bytes (small chunk) are denominated Small bins.**
 - Chunks in small bins are called small chunks.
 - A Small bin will contain small chunks of all the same size.
 - Spaced 8 bytes - 32bit
 - Spaced 16 bytes - 64 bit
 - Exact fit search.
 - Placed from index 2 - 64 in bins.
- **Bins for chunks sizes >= 512 || 1024 (large chunk) bytes are denominated Large bins.**
 - Large bins contain chunks called large chunks.
 - Are approximately logarithmically spaced.
 - They will contain chunks within a particular size range.
 - Best fit search.
 - Placed from index 65 - 127 in bins.

Recap:

- **bins[2] to bins[64]** - Small bins.
- **bins[65] to bins[127]** - Large bins.

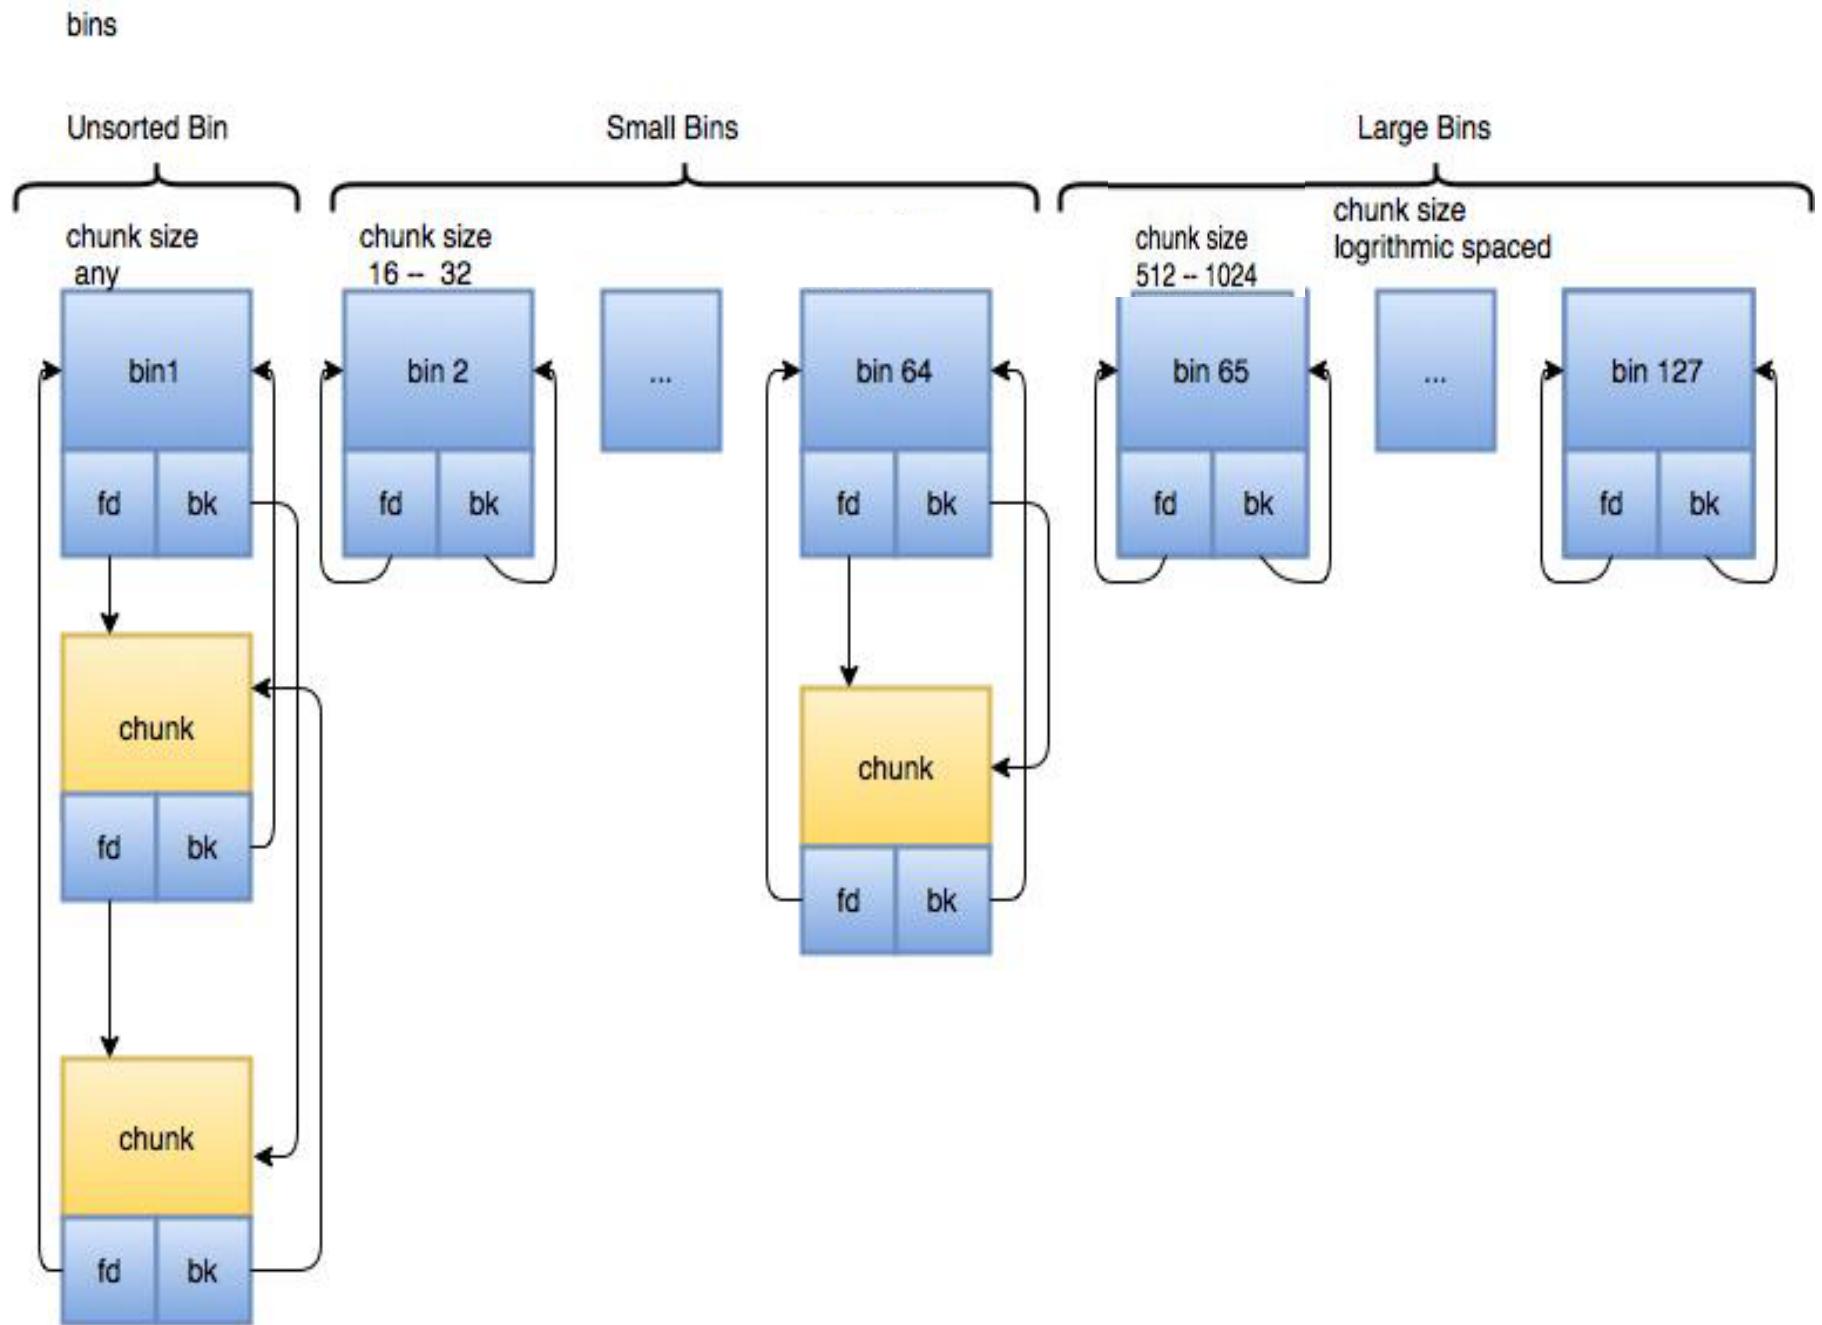
What about bins[0] and bins[1]?

- **bins[0] does not exist.** (malloc.c #1459)
- **bins[1] is the Unsorted bin.**

Unsorted bin:

- Intended to reuse recent freed non fastchunks of not particular size.
- For all deallocations:
 - If chunk is a fastchunk: once freed is placed in a fastbin.
 - Else: once freed chunk is first placed in the unsorted bin.
- Chunks in the Unsorted bin are not sorted by chunk size as in other bins.
- Malloc will sort the unsorted bin in chunk requests in which:
 - Malloc did not find an exact fit in fastbins or small bins for demanded chunk.
 - Malloc can't consolidate any fastchunk to return the user a suitable chunk from demand.
(more details in coming slides).

Heap Structures

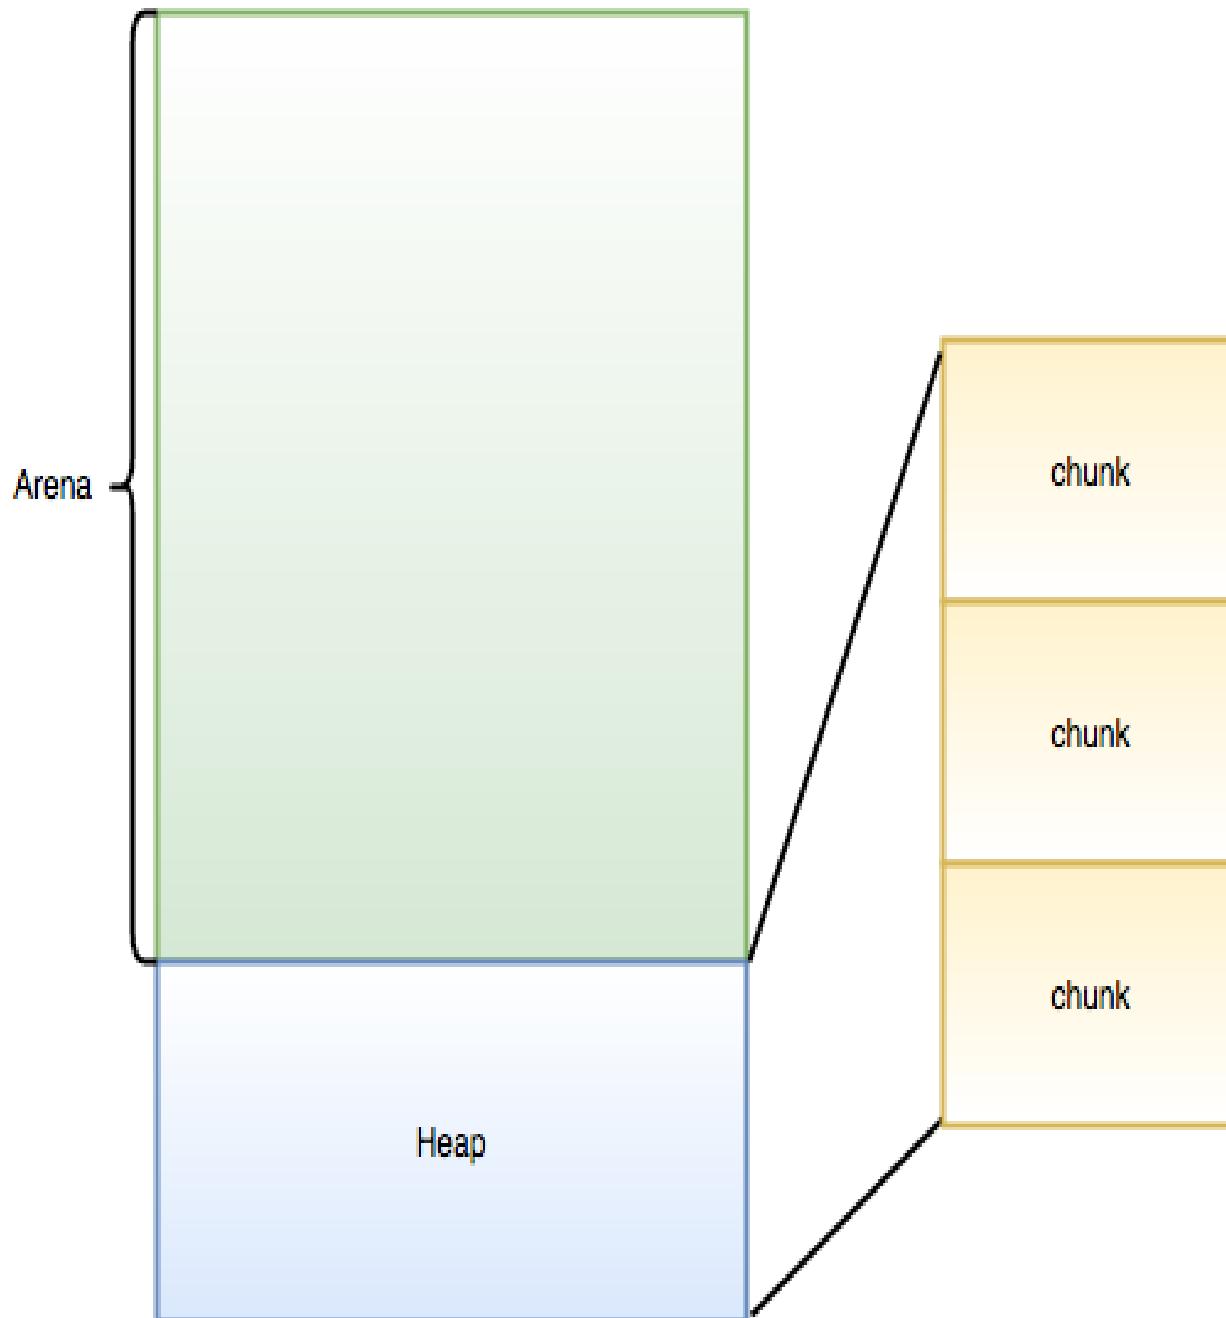


Heap Structures

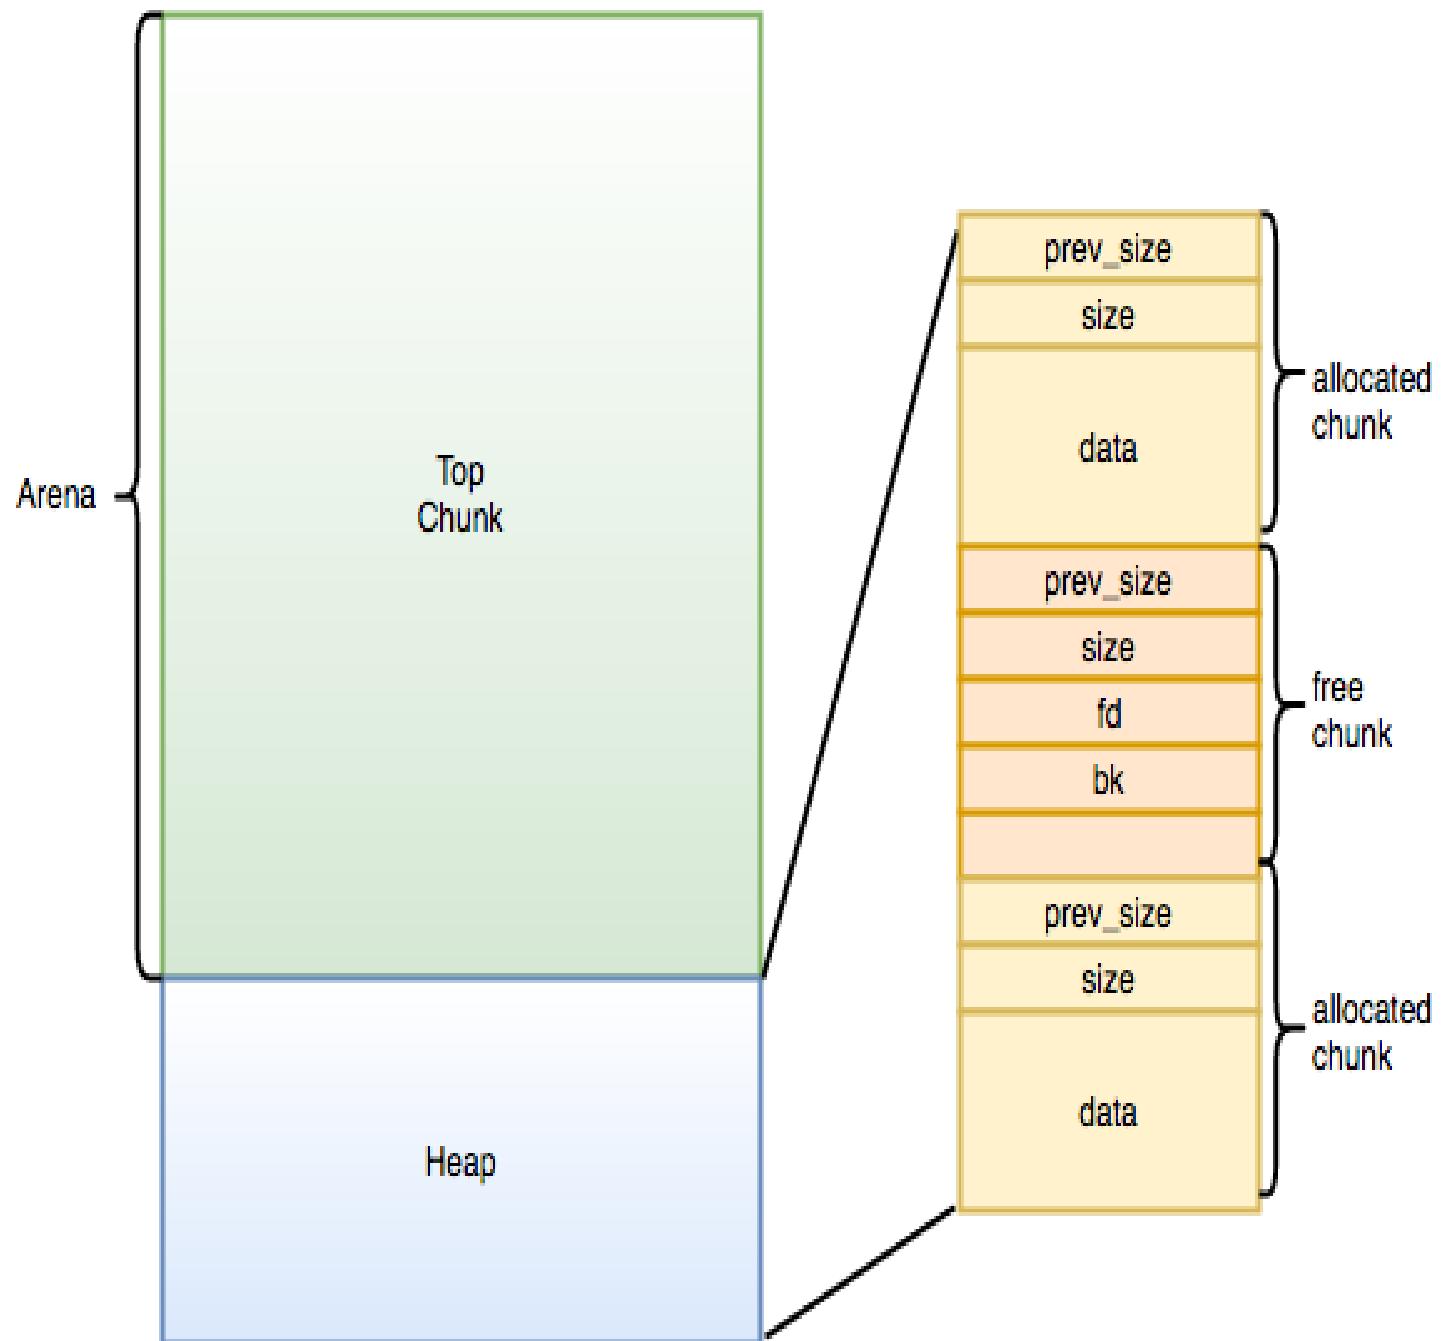


- If ideal chunk for malloc demand is not found in Free Lists, a chunk will need to be retrieved from the thread's arena in order to satify malloc's demand.
- How is an arena actually implemented?
 - The arena of a thread is implemented in the form of a chunk.
 - This chunk receives the name of **top chunk**.
 - It is never included in any bin.
 - is used only if no other chunk is available.
 - Top chunk is independent of heaps.

Conceptual model of Arena, Heap and Chunk



Practical model of Arena, Heap and Chunk

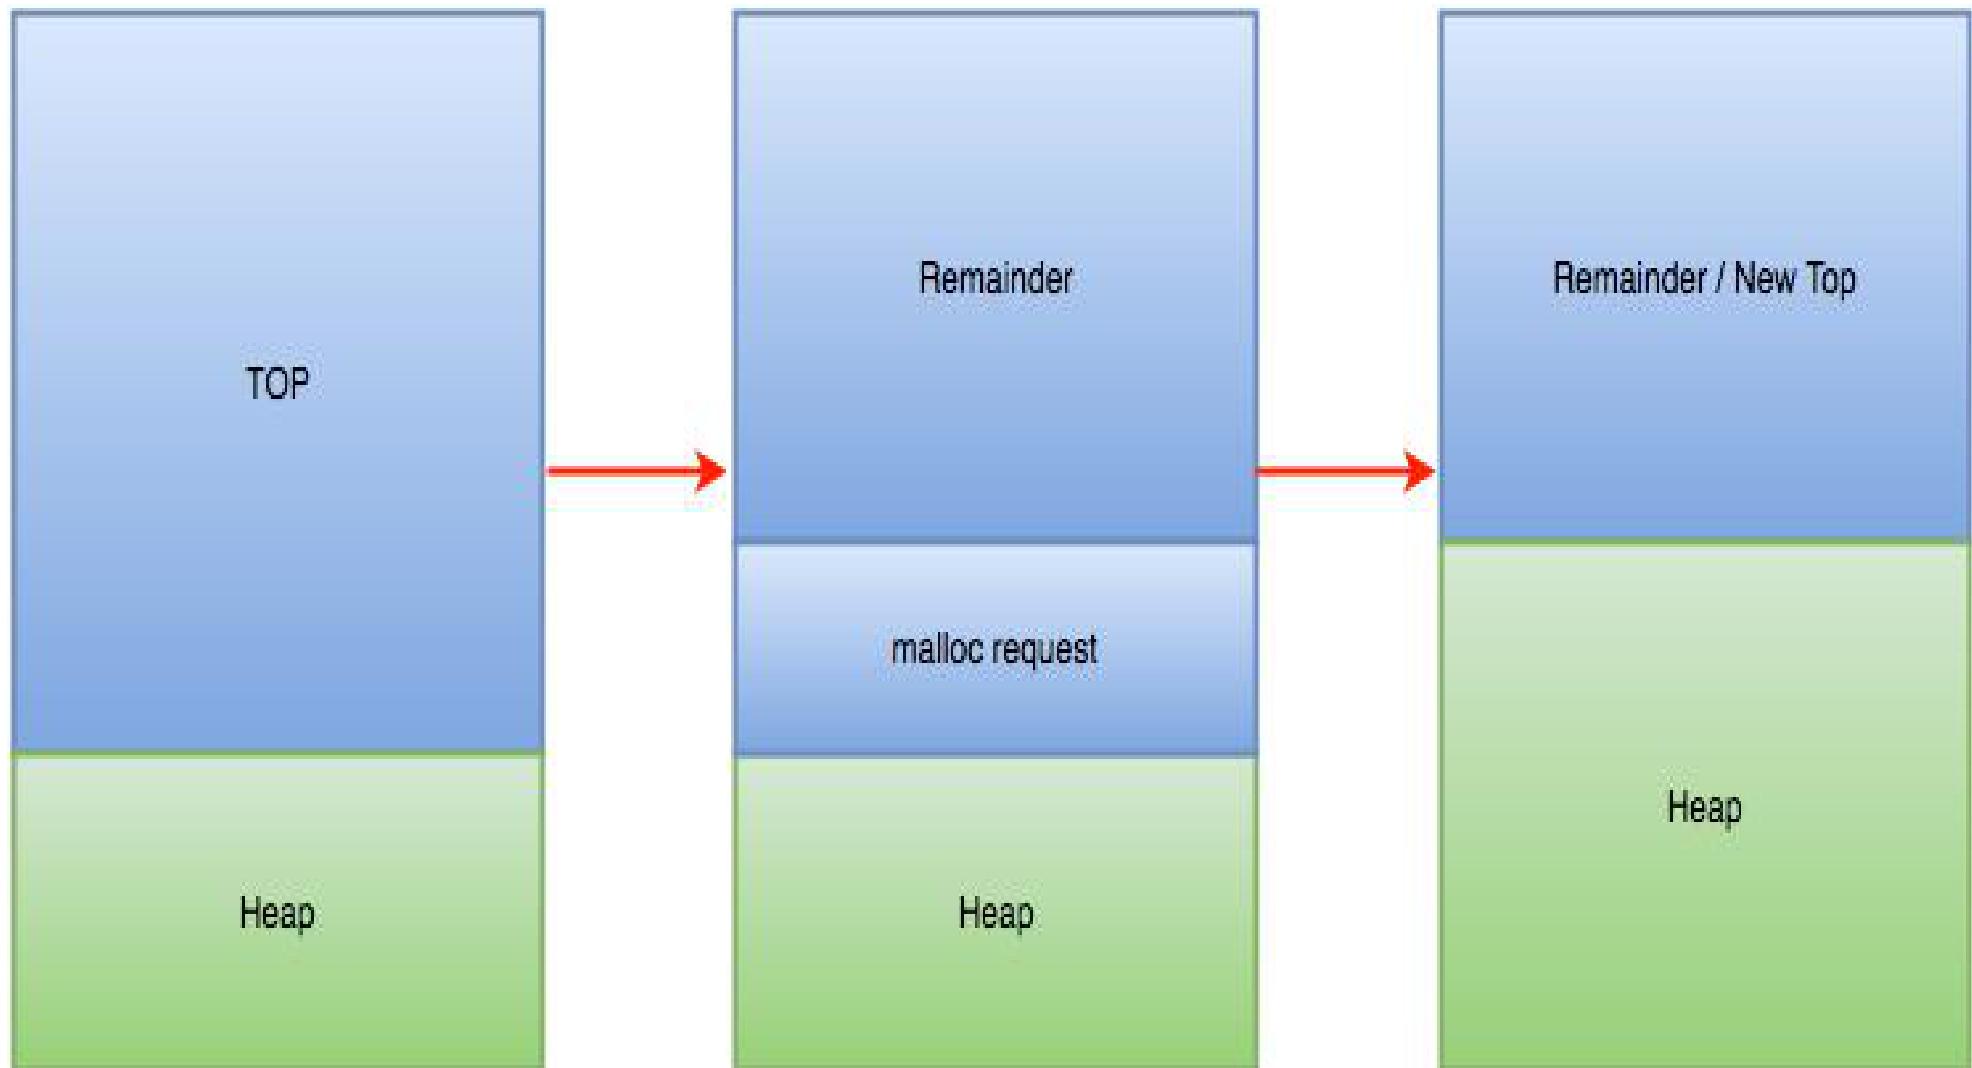


Heap Structures



- **The remainder chunk.**
 - Product of retrieving a chunk from arena
 - If ideal chunk for malloc demand was not found in free lists, memory will need to be retrieved from the thread's arena.
 - top chunk has to be splitted.
 - splitted into two parts:
 - Malloc requested size chunk
 - Remainder chunk (new top).

Retrieving a chunk from the heap segments's top chunk



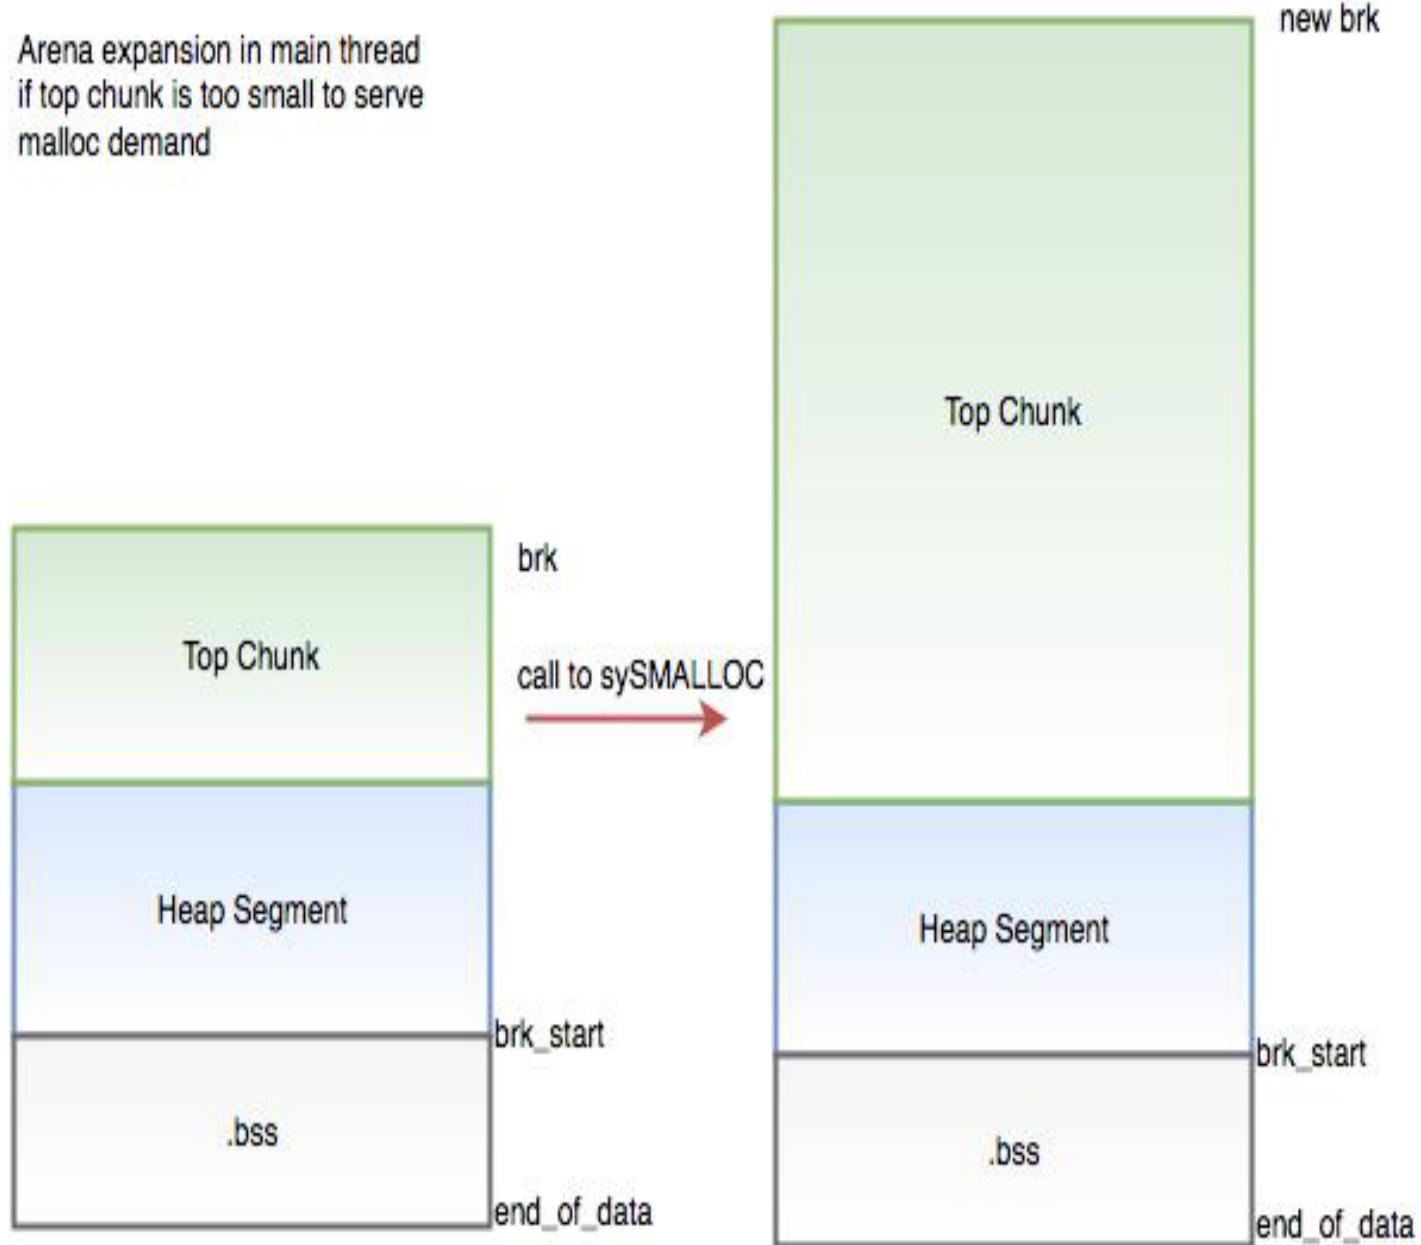
Heap Structures



- **What about if the thread's arena is not big enough to serve a particular malloc demand?**
 - Thread's arena will need to be expanded.
 - This process is different in child threads than it is for the main thread.
 - For the main thread:
 - Arena can be expanded by calling `sbrk()`, therefore expanding the dimensions of the heap segment by incrementing program break.

Main thread's arena expansion

Arena expansion in main thread
if top chunk is too small to serve
malloc demand



Heap Structures



- **For child threads:**
 - **New mmap() call will need to be made, meaning:**
 - **New arena will be created and combined with previous arena.**
 - **New top chunk will be generated.**
 - **Previous top chunk will become a free chunk.**
 - **New Heap will be created in the newer mmapped region context.**
 - **These two mmapped regions will become part of the same arena, however they most likely will not be adjacent in memory.**
 - **To manage multiple heaps in child threads an additional structure is used. This structure is `heap_info`.**

Heap Structures



- **heap_info**

```
typedef struct _heap_info {
    mstate ar_ptr;           /* Arena for this heap. */
    struct _heap_info *prev;  /* Previous heap. */
    size_t size;              /* Current size in bytes. */
    size_t mprotect_size;     /* Size in bytes that has
                                been
                                PROT_READ|
                                PROT_WRITE. */
    char pad[-6 * SIZE_SZ & MALLOC_ALIGN_MASK];
//alignment
} heap_info;
```

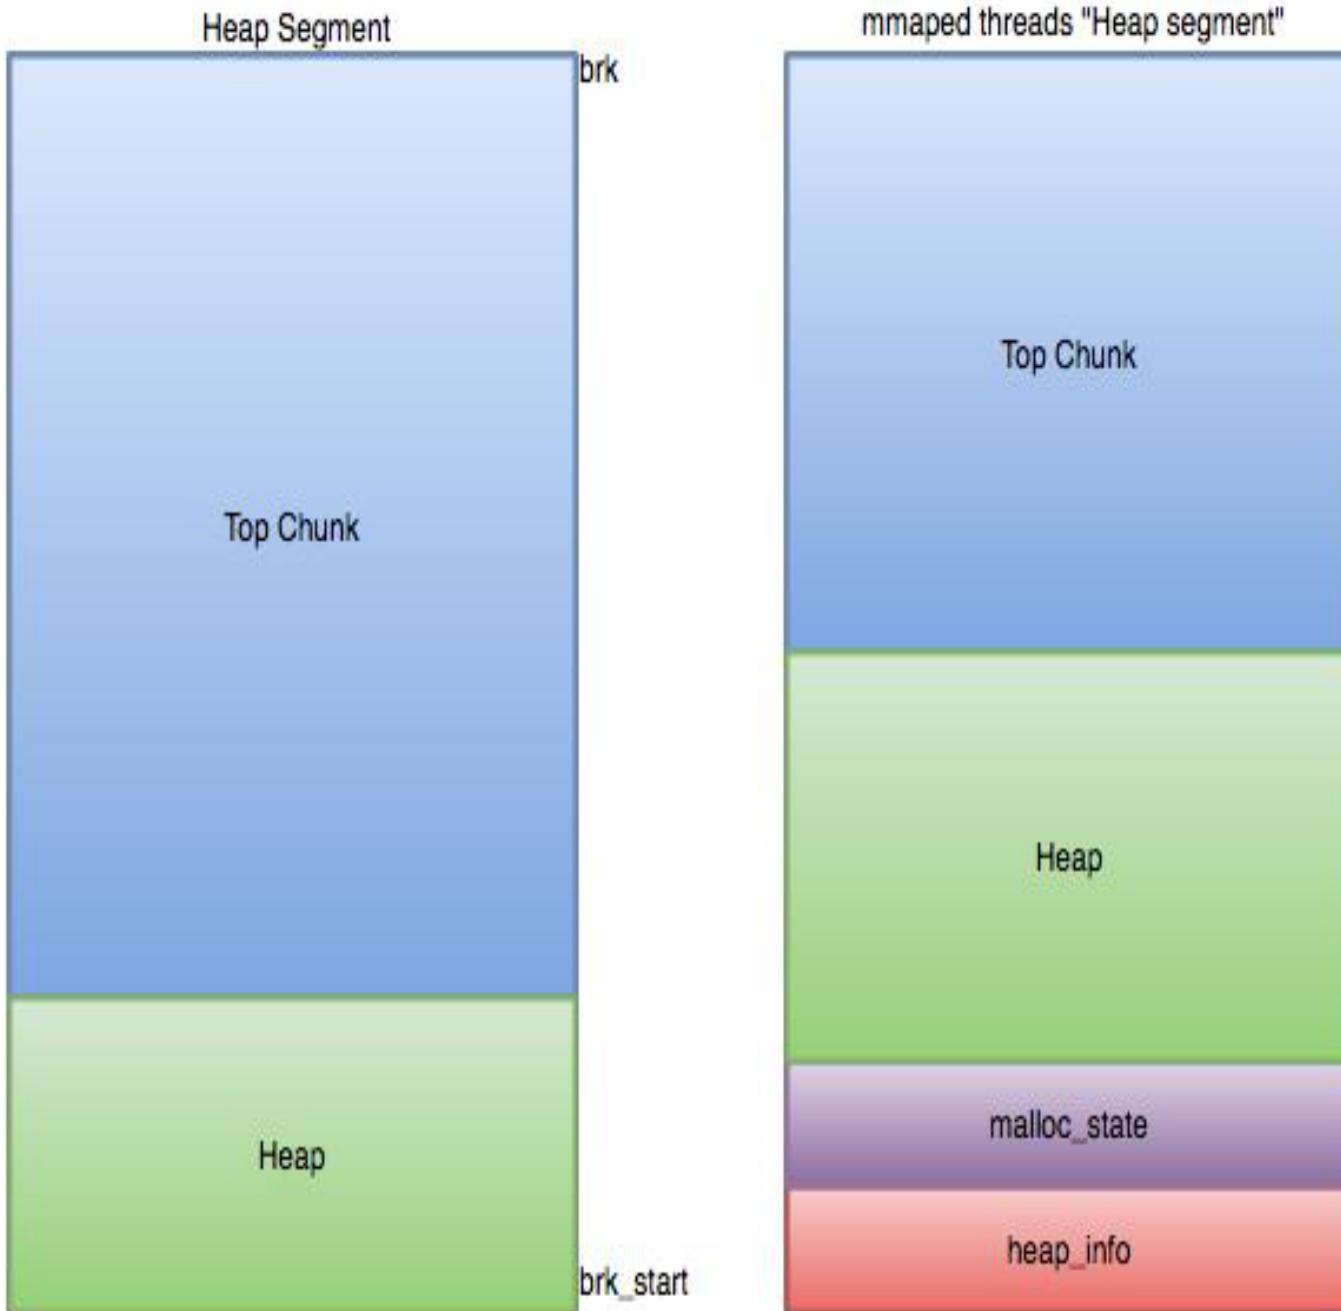
Heap Structures



- **heap_info**

```
mstate ar_ptr;          /* Arena for this heap. */  
struct _heap_info *prev; /* Previous heap. */  
size_t size;            /* Current size in bytes. */
```

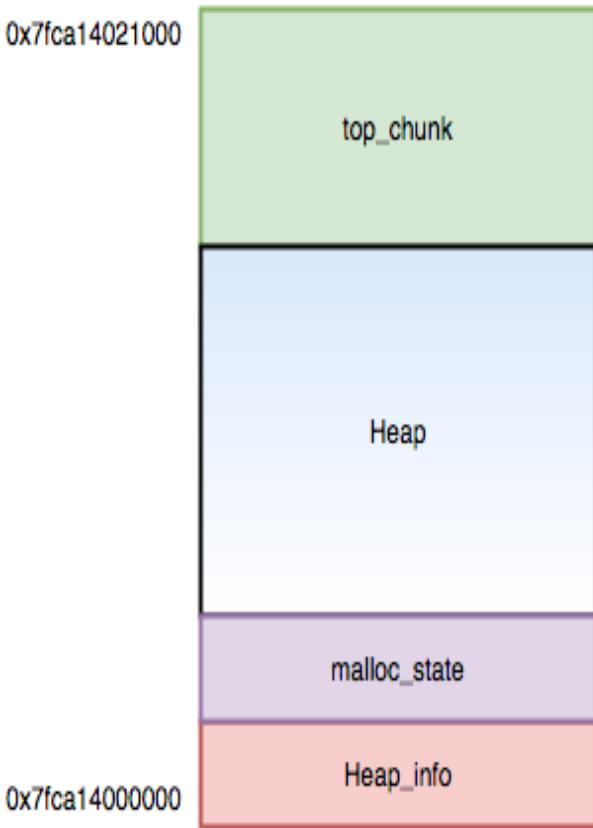
- Intended to keep track of multiple heaps in child threads.
- Multiple heaps are not supported in the main thread.
 - `heap_info` does not exist in single threaded applications.
- The main thread also uses a different instance of `malloc_state` (`main_arena`).
- Instances for `malloc_state` and `malloc_info` in child threads are located at the beginning of the their particular arenas.



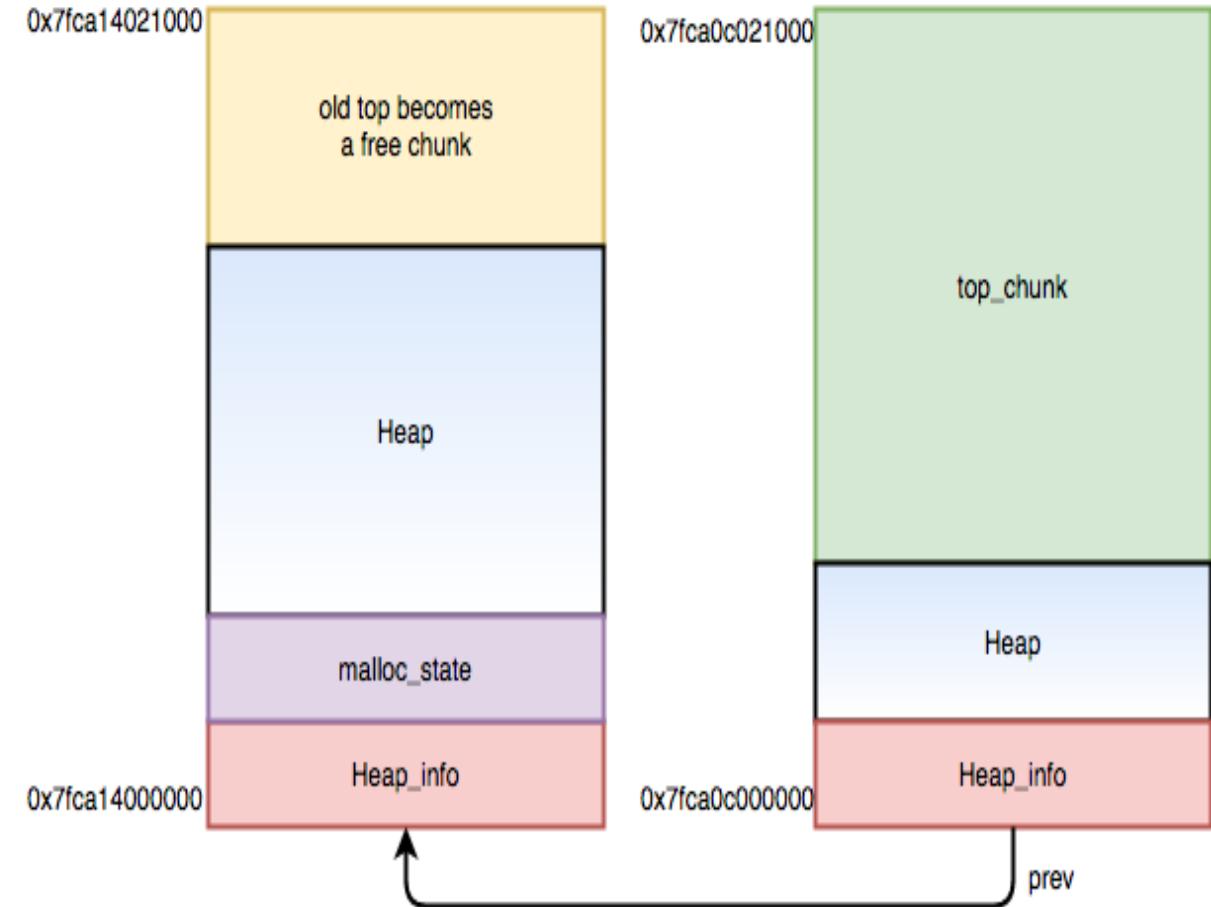
Heap Structures



Thread arena Before new mmap() call

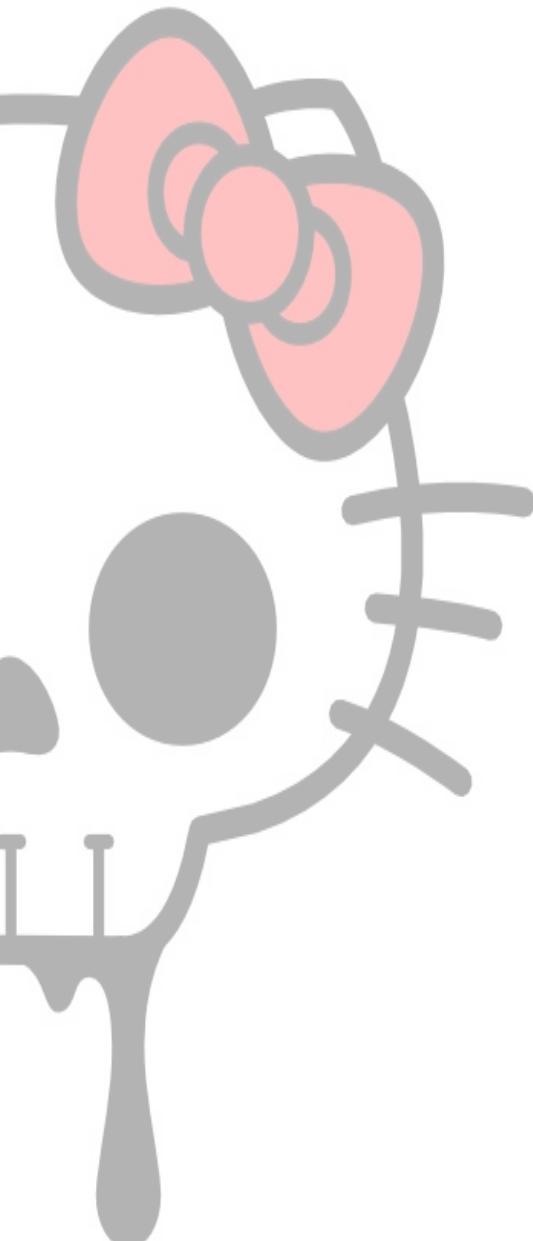


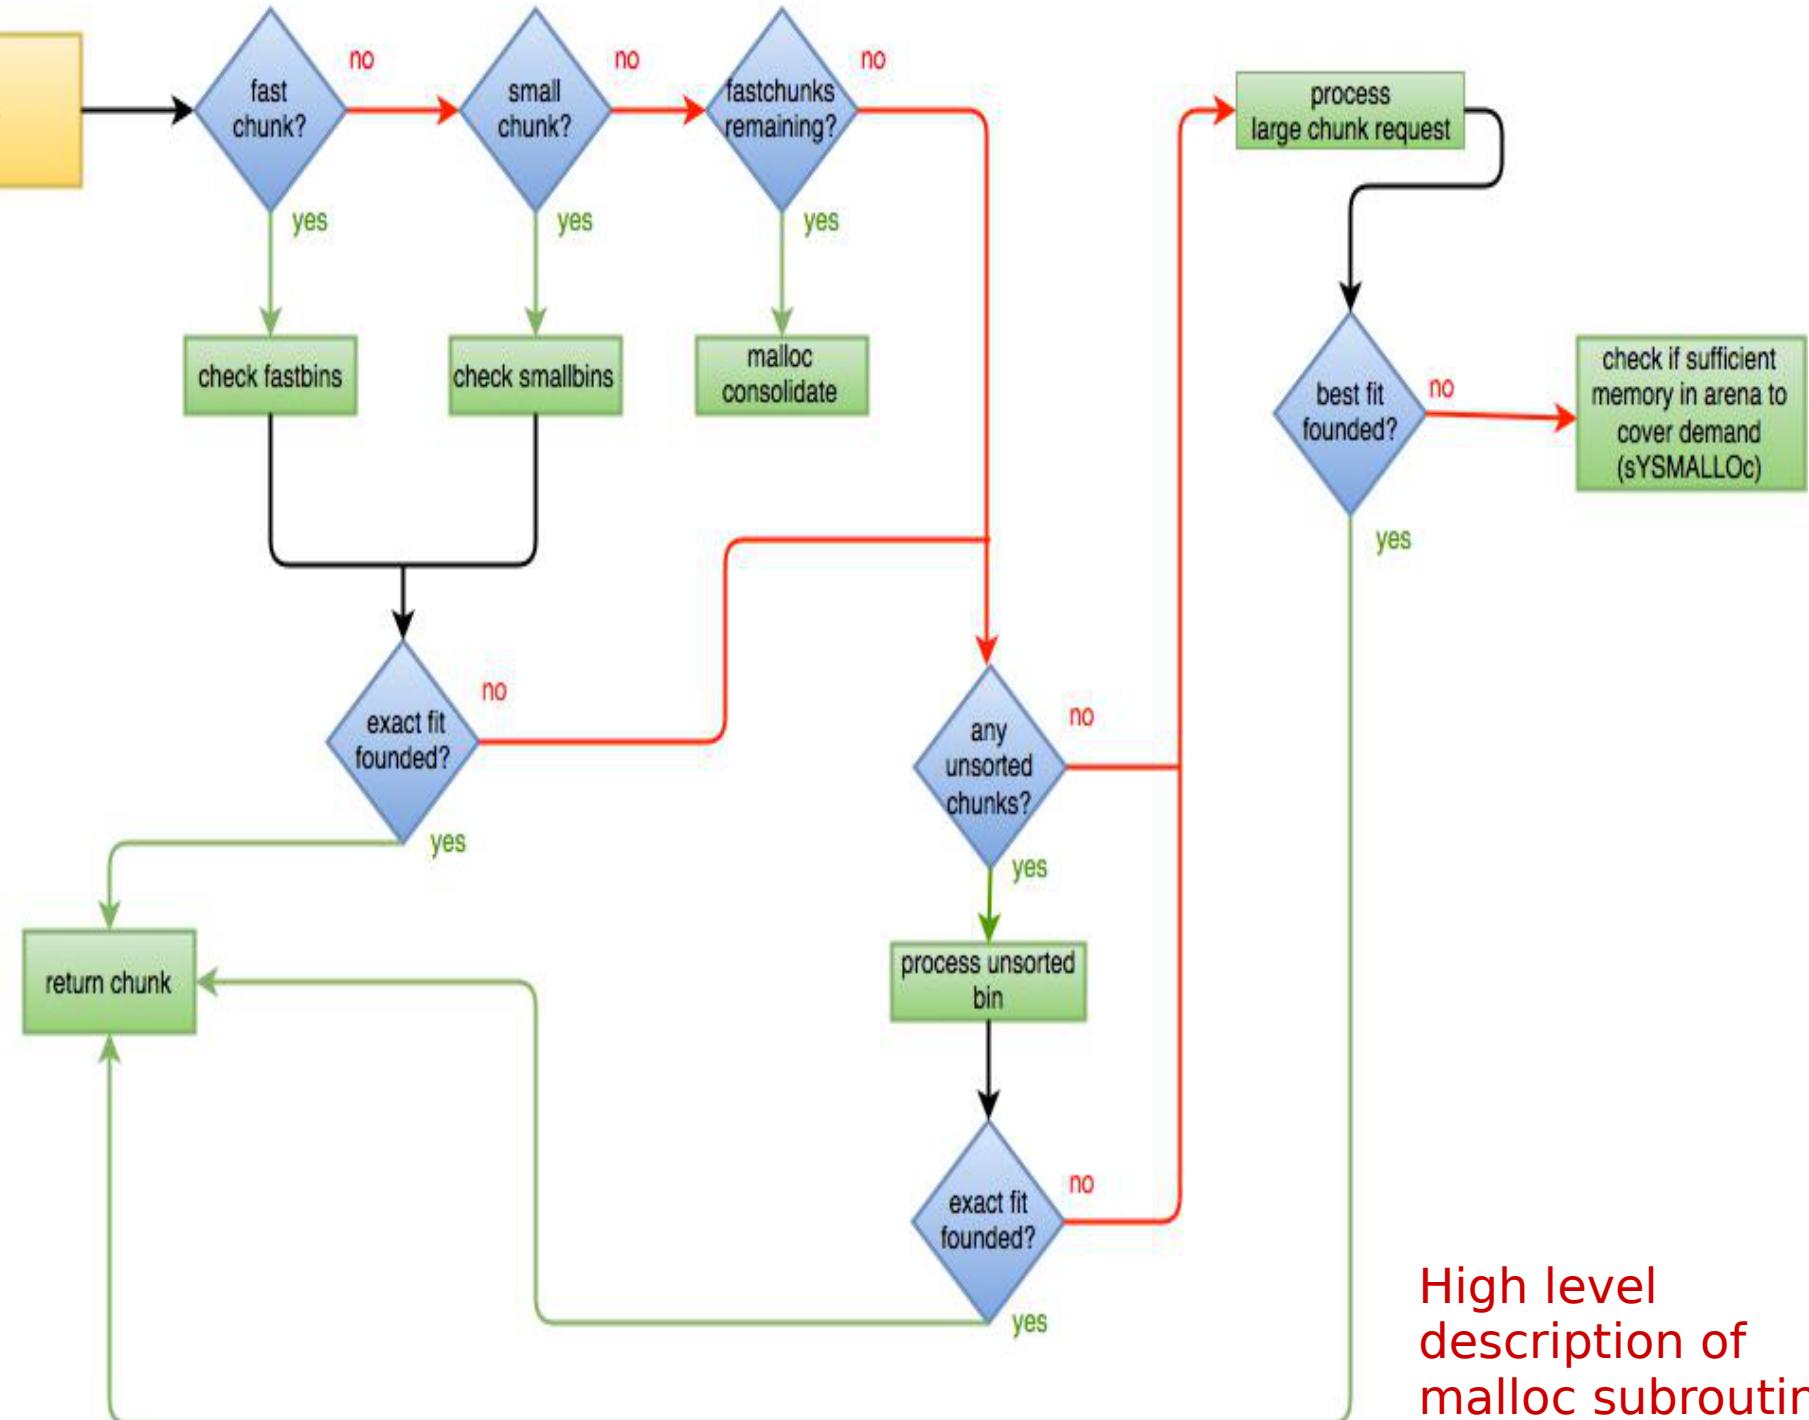
Thread arena after new mmap() call



Introduction to glibc malloc algorithm

amnesia



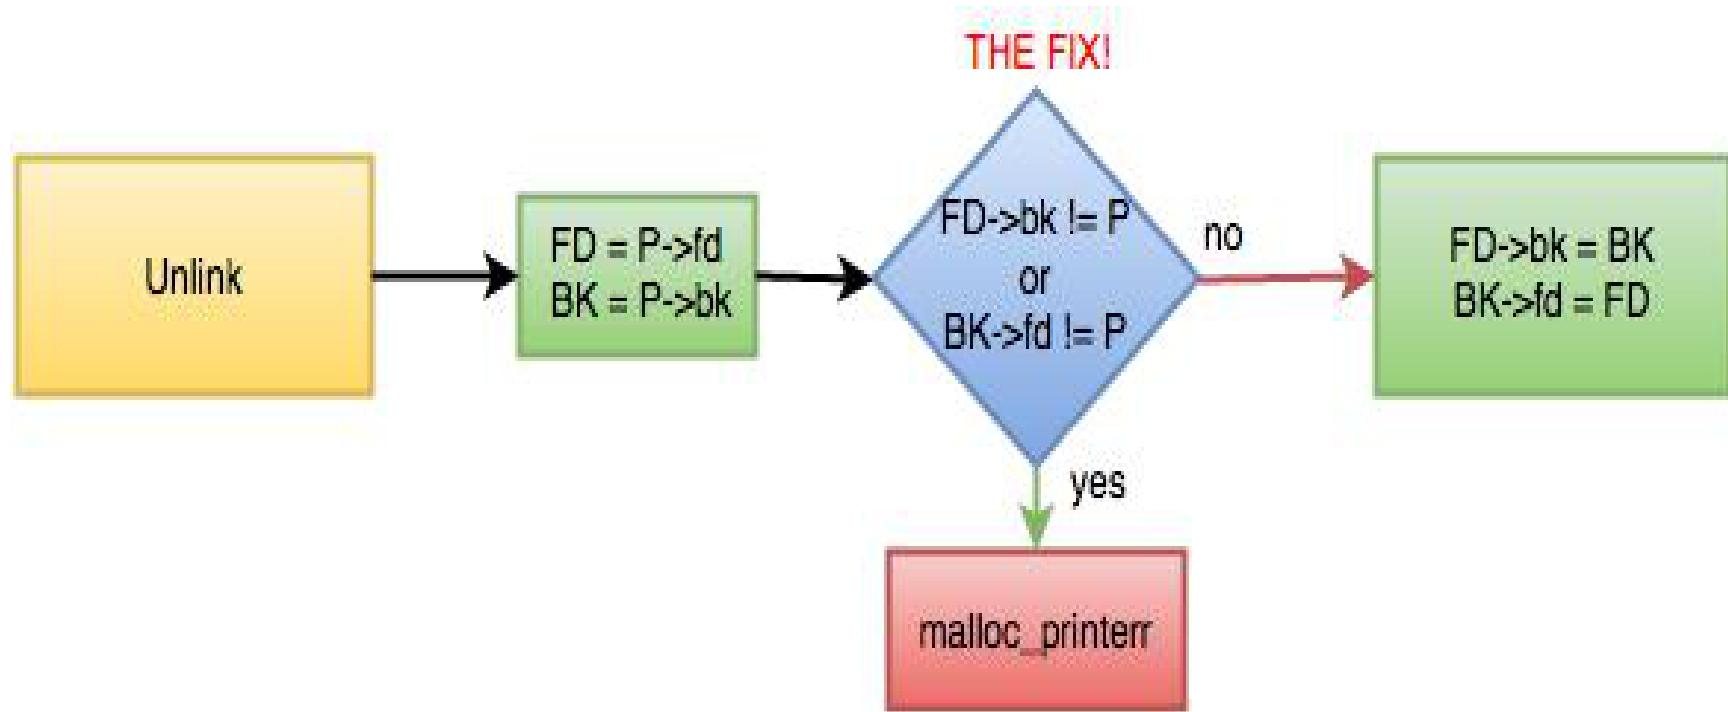


High level
description of
malloc subroutine

The unlink macro:

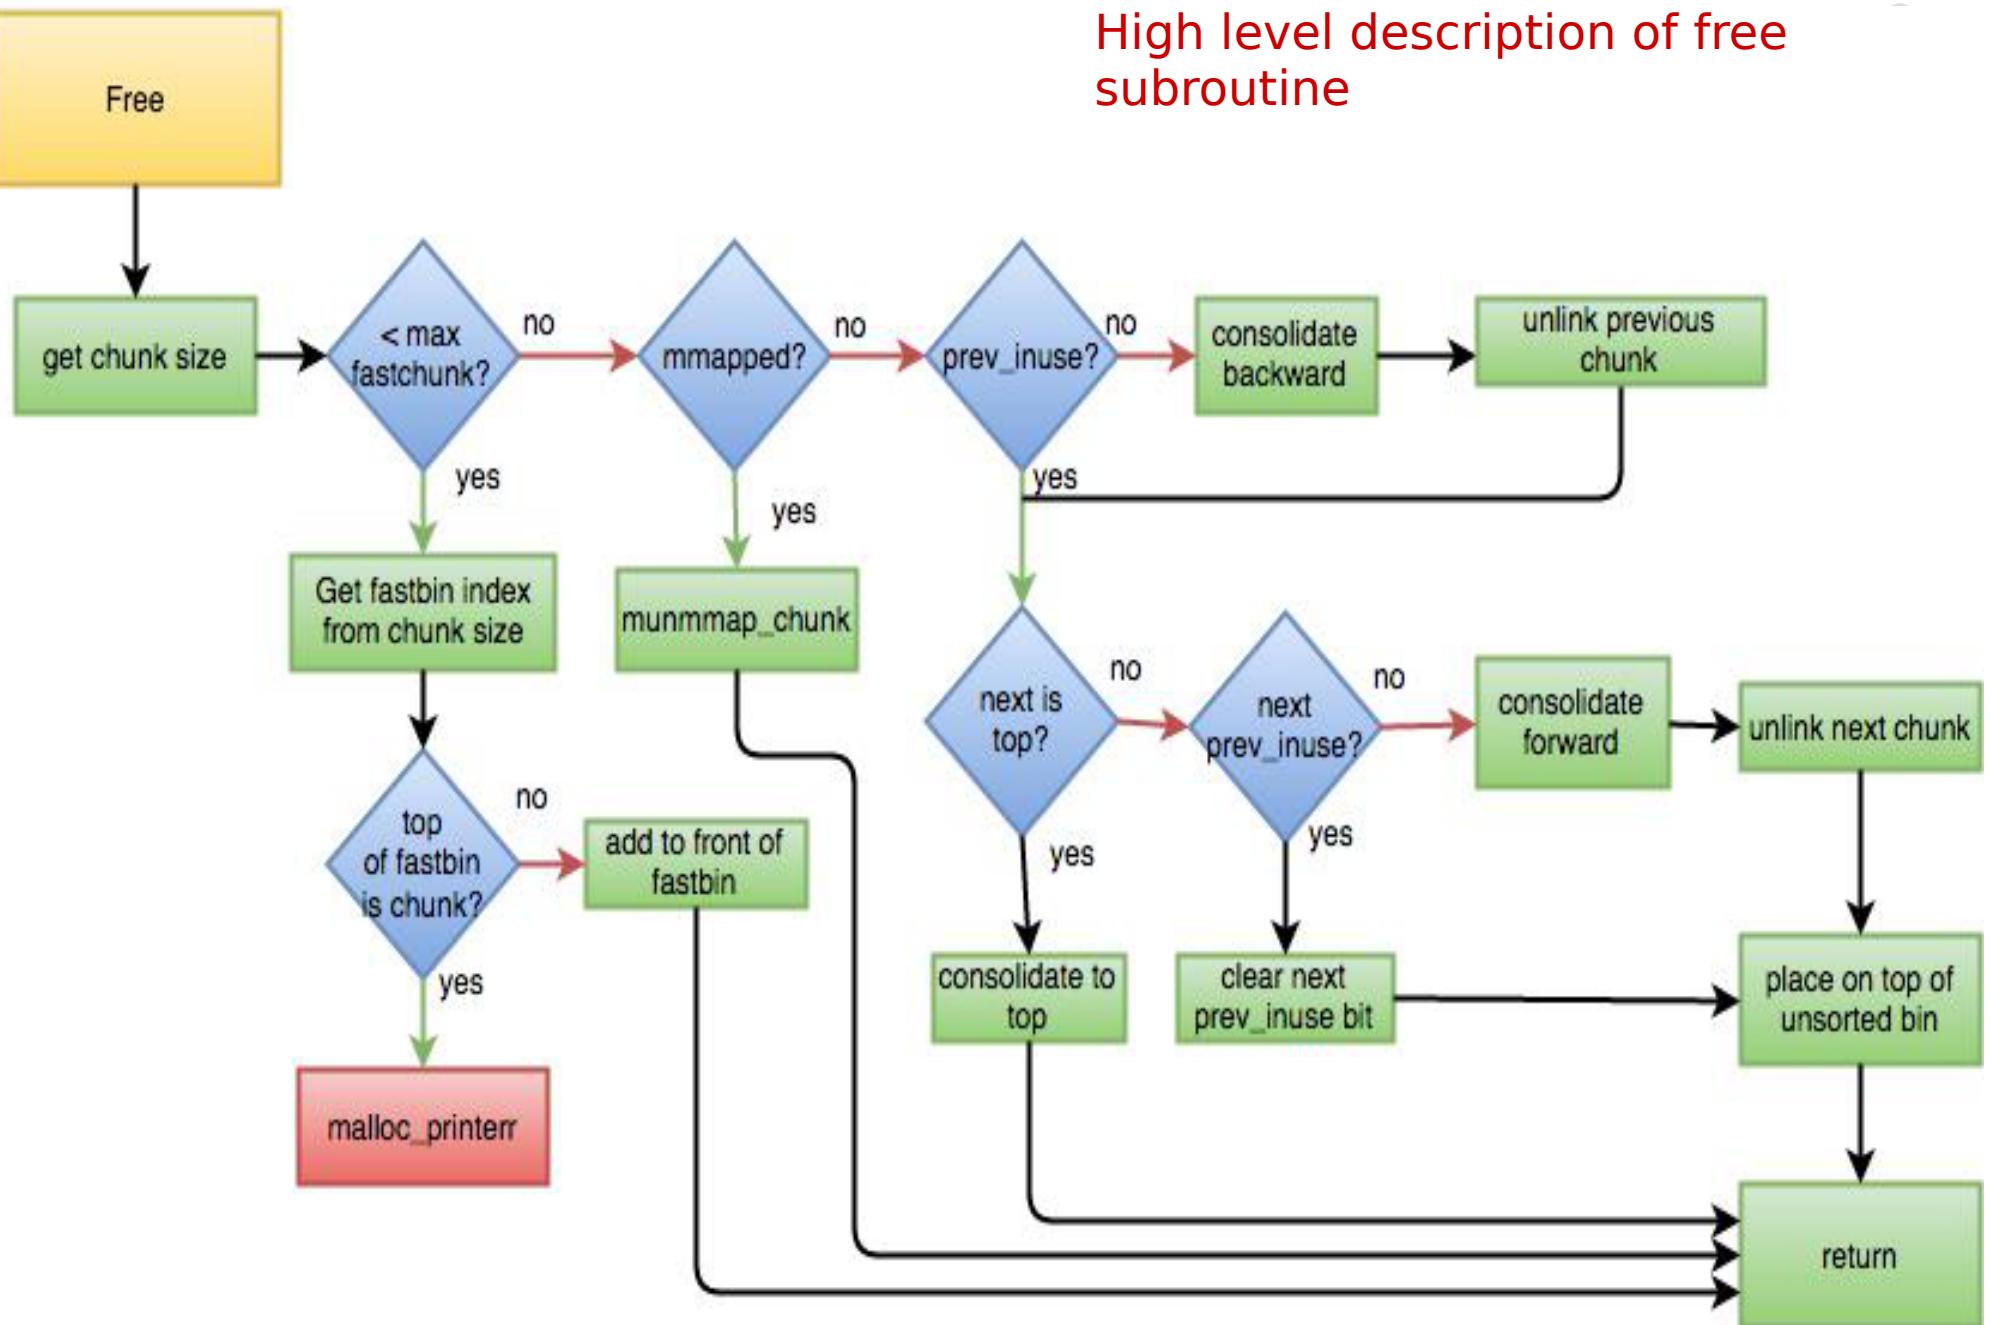
- Intended to take a chunk off a free list.
- Highly exploited in previous version of glibc malloc.
- Fixed (not really).

Introduction to glibc malloc algorithm



High level description of unlink macro

High level description of free subroutine



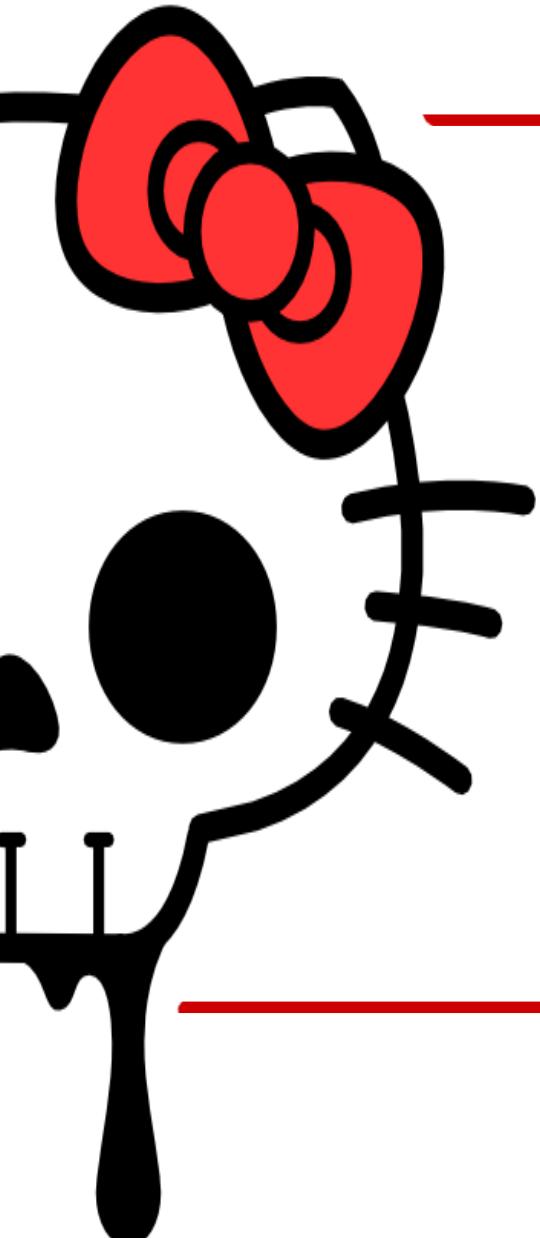
Dmh command family

```
[0x0040057d]> dmh?  
Usage: dmh # Memory map heap  
dmh                                         List chunks in heap segment  
dmh [malloc_state]                           List heap chunks of a particular arena  
dmha                                         List all malloc_state instances in application  
dmhb                                         Display all parsed double linked lists of main_arena's bins instance  
dmhb [bin_num|bin_num:malloc_state]          Display parsed double linked list of bins instance from a particular arena  
dmhbg [bin_num]                             Display double linked list graph of main_arena's bins instance [Under development]  
dmhc @[chunk_addr]                          Display malloc_chunk struct for a given malloc chunk  
dmhf                                         Display all parsed single linked lists of main_arena's fastbinY instance  
dmhf [fastbin_num|fastbin_num:malloc_state]  Display single linked list in fastbinY instance from a particular arena  
dmhg                                         Display heap graph of heap segment  
dmhg [malloc_state]                          Display heap graph of a particular arena  
dmhi @[malloc_state]                         Display heap_info structure/structures for a given arena  
dmhm                                         List all elements main thread's malloc_state struct(main_arena)  
dmhm [malloc_state]                          List all elements for a given malloc_state instance  
dmh?                                         Show map heap help  
[0x0040057d]> ■
```

Dmh command family

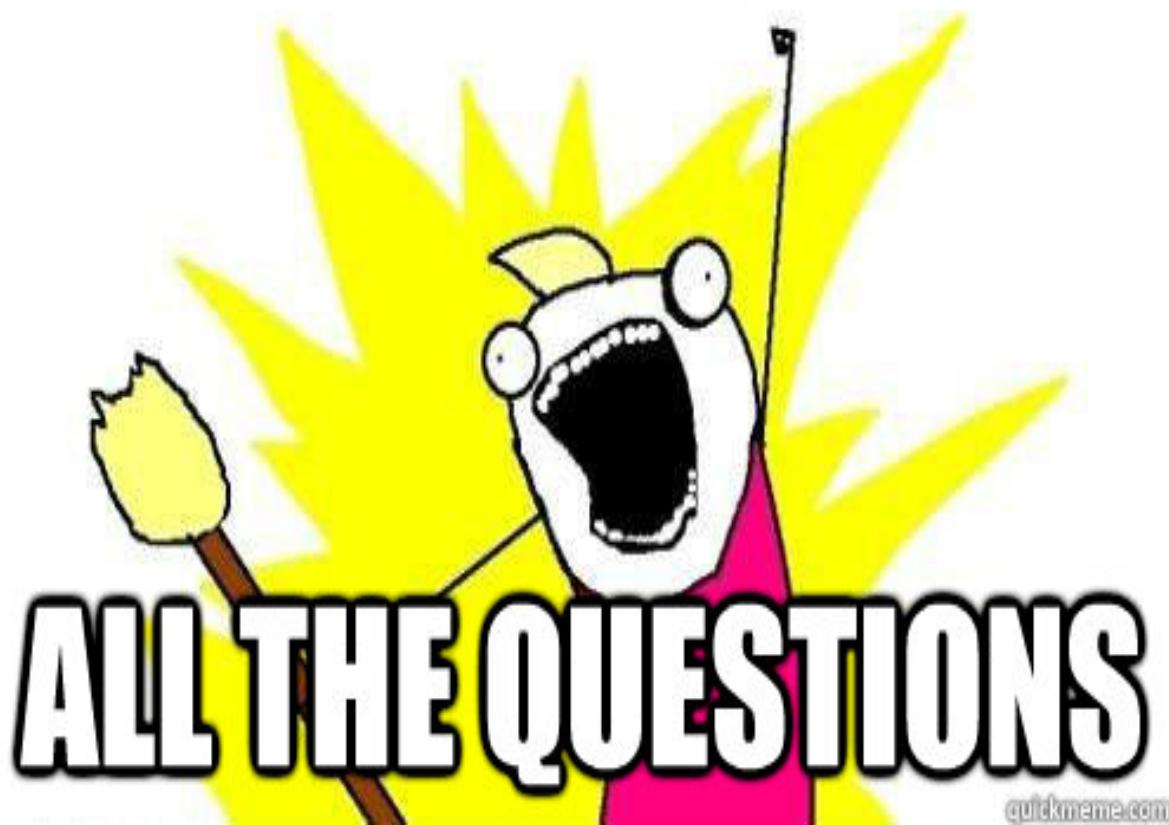


Conclusion

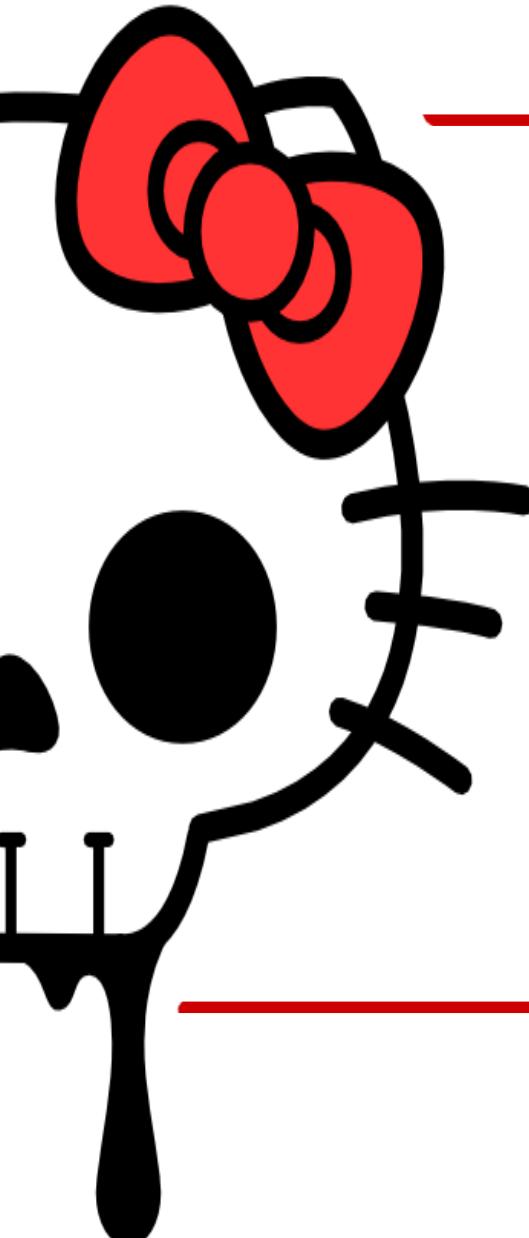
- 
- The extension does give a more convenient view of the heap.
 - Overall experience was satisfactory.
 - Learned to defend ourselves with Radare2 API
 - Learned in great extend how the glibc heap works.
 - Extension has not reach its maximum potential.
 - Not all strucutures have been implemented yet (malloc_par).
 - Planning to do similar extensions for different allocators such as:
 - Magazine.
 - Jemalloc.

Questions?

ASK



References

- 
- **Glibc malloc source code:**
 - https://fossies.org/dox/glibc-2.24/malloc_8c_source.html.
 - **Relevant documentation:**
 - <https://exploitfun.wordpress.com/2015/02/10/understanding-glibc-malloc/>
 - <https://sourceware.org/glibc/wiki/MallocInternals>
 - <http://files.cnblogs.com/files/hseagle/demo.pdf>
 - <http://www.cs.cmu.edu/afs/cs/academic/class/15-213-f10/www/lectures/17-allocation-basic.pdf>